

Message Passing Concurrency

BY RICCARDO TERRELL - @TRIKACE

Agenda

- Agents motivations
- What are Agents
- Agents vs Actors
- Agents patterns

Agent motivations

- You can manage **shared data** and resources **without locks**.
- You can easily follow the **SRP**, because each **agent** can be designed to **do only one thing**.
- It encourages a "pipeline" model of programming with "producers" sending messages to decoupled "consumers"
- It is straightforward to **scale**
- Errors can be handled gracefully, because the decoupling means that agents can be created and destroyed without affecting their clients.

Agent-based concurrency

Programs **composed** from agents

- Agents can be viewed as “running” objects

Agents **exchange** messages

- Receive message and react
- Trigger event when work is done

Reactive system

- Handle inputs while running
- Emit results while running

Agent is a single and independent unit of computation



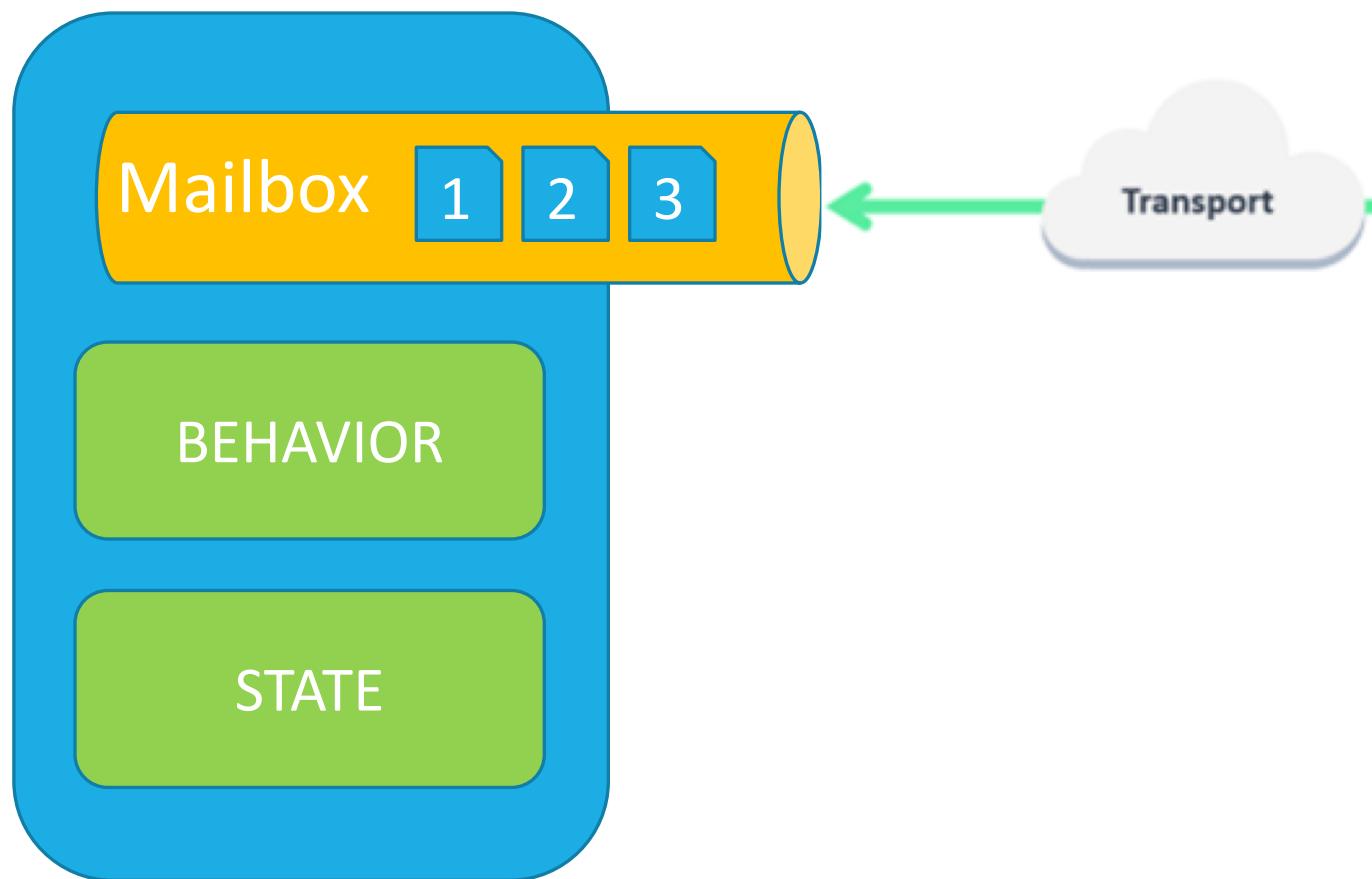
Message-Driven



What are Agents

- Communication between agents done by **sending message**
- The messages are put on a **queue** for each agent
- Each agent process **one message** at time

Message Passing based concurrency



- Processing
- Storage – State
- Communication only by messages
- Share Nothing
- Message are passed by value
- Lightweight object
- Running on it's own thread
- No shared state
- Messages are kept in mailbox and processed in order
- Massively scalable and lightening fast because of the small call stack

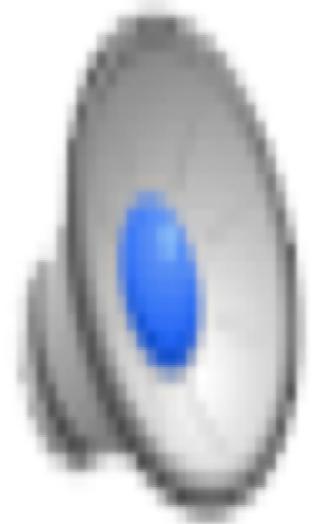
Share Nothing, Async Message Passing

- » Encourages shared-nothing process abstractions
 - mutable state - private
 - shared state - immutable
- » Asynchronous message passing
- » Pattern matching

Declaring messages

- » Easier to reason about
- » Higher abstraction level
- » Easier to avoid
 - Race conditions
 - Deadlocks
 - Starvation
 - Live locks
- » Distributed computing

What is an Actor?



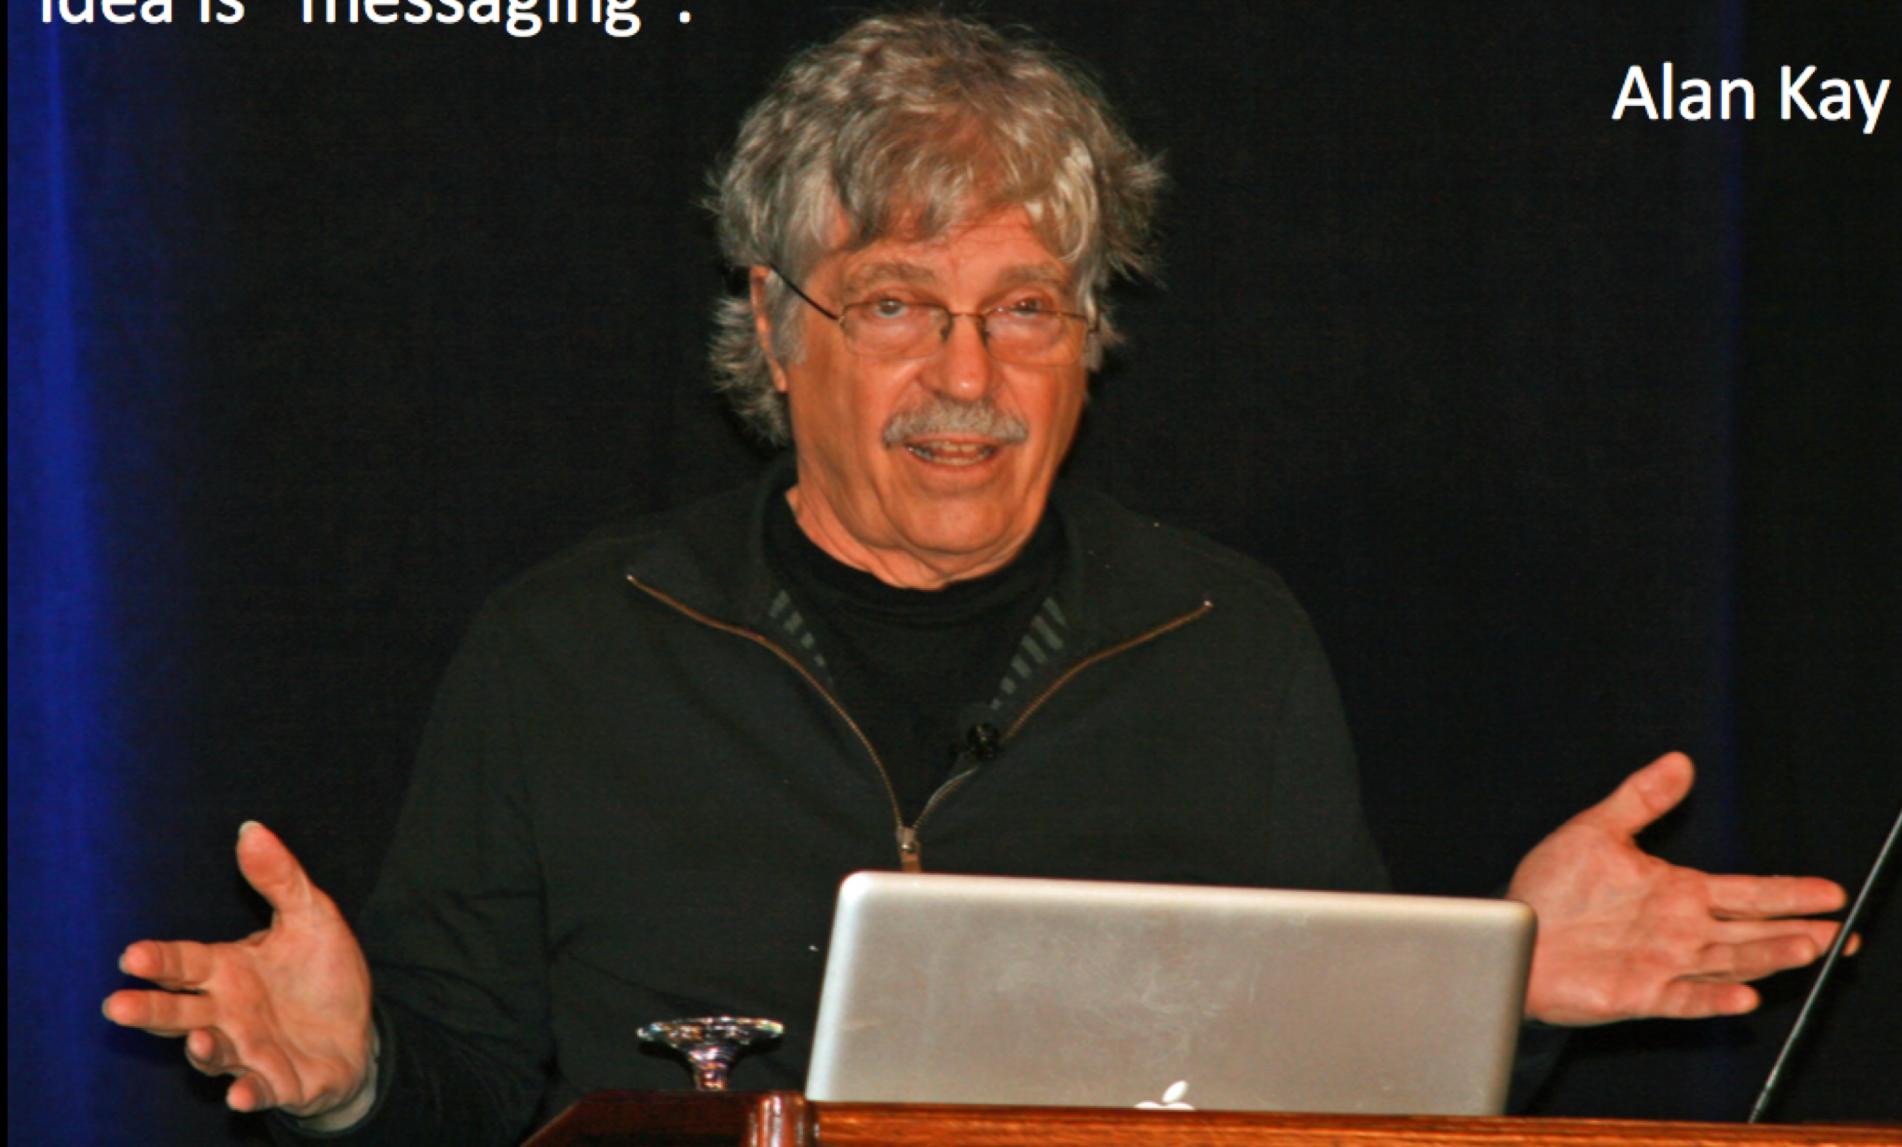
- **Share Nothing**
- **Message are passed by value**
- Light weight processes/threads communicating through messaging
- Communication only by messages
- Lightweight object
- Processing
- Storage – State
- Running on it's own thread.
- Messages are buffered in a “mailbox”

Agents vs Actors

- Agent work in process
- Actors work also out of process
- Agent messages pass by reference
- Actor messages always serialized before sent

I'm sorry that I coined the term "objects", because it gets many people to focus on the lesser idea. The big idea is "messaging".

Alan Kay



Simple agent in F#

Receive message and say “Hello”

```
let hello = Agent.Start(fun agent -> async {
    while true do
        let! name = agent.Receive()
        printfn "Hello %s" name
        do! Async.Sleep 500 })

hello.Post("World!")
```

- Single instance of the body is running
- Waiting for message is asynchronous
- Messages are queued by the agent

Mutable and immutable state

Mutable state

- Accessed from the body
- Used in loops or recursion
- Mutable variables (ref)
- Fast mutable collections

Immutable state

- Passed as an argument
- Using recursion (**return!**)
- Immutable types
- Can be returned from the agent

```
Agent.Start(fun agent -> async {
    let names = ResizeArray<_>()
    while true do
        let! name = agent.Receive()
        names.Add(name) })
```

```
Agent.Start(fun agent ->
    let rec loop names = async {
        let! name = agent.Receive()
        return! loop (name::names) }
    loop [])
```

Declaring messages

Agents handle multiple messages

- Message type using discriminated union

```
type CacheMessage<'T> =
| Add of string * 'T
| Get of string * AsyncReplyChannel<option<'T>>
| Reset
```

Safety guarantees

- Agent will be able to handle all messages

Agent-based programming

Program consists of agents

- **Lightweight** – can create lots of them
- Written using **asynchronous** workflows

Agents communicate via messages

- Communication is **thread-safe**
- Messages are **queued**

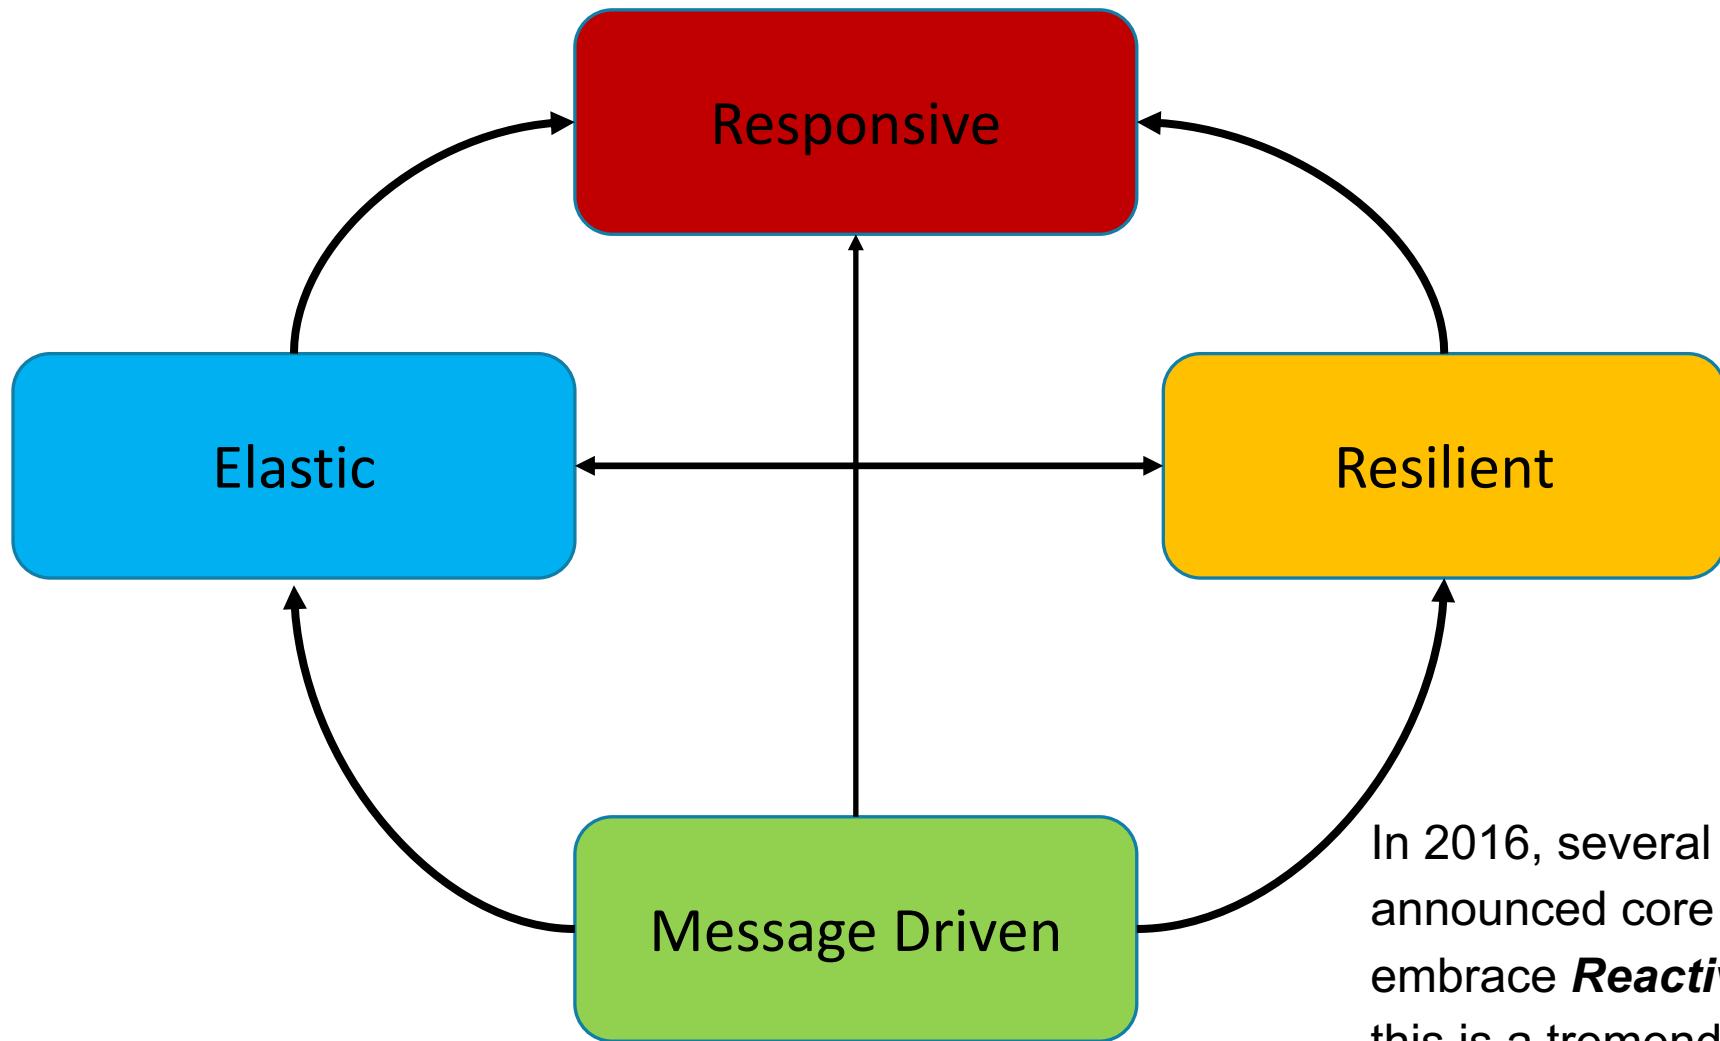
Enables parallelism

- **Different agents vs. multiple instances**

Actors / Agents

- » React to received messages by executing a behavior function
 - can only change the state of the actor itself
 - can send messages to other actors
- » Actors never share state and thus never need to compete for locks for access to shared data
- » Actors exchange data by sending **immutable** messages
- » Messages are sent asynchronously
- » Actors do not block waiting for responses to their messages

Reactive Manifesto



In 2016, several major vendors have announced core initiatives to embrace **Reactive Programming** this is a tremendous validation of the problems faced by companies today.

Reactive Manifesto & Actor Model

Responsive

Event Driven

Message-Driven

Communication by messages

Resilient

Fault tolerant by Supervision

Elastic

Clustering and Remoting across multiple machines

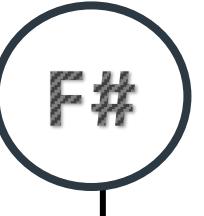
Immutability OR Isolation

```
let printerAgent = MailboxProcessor.Start(fun inbox->
    let list = new List<string>()
    // the message processing function
    let rec messageLoop() = async{
        // read a message
        let! msg = inbox.Receive()
        // process a message
        printfn "message is: %s" msg
        // loop to top
        list.Add(msg)
        return! messageLoop()
    }
    // start the loop
    messageLoop()
)
```

F#

Send message to agent

```
printerAgent.Post "hello"  
printerAgent.Post "hello again"  
printerAgent.Post "hello a third time"
```

A circular icon containing the F# logo, which consists of the letters 'F' and '#' stacked vertically.

F#

Agent-based architecture

Lots of things going on!

- How to keep a big picture?

Using loosely coupled connections

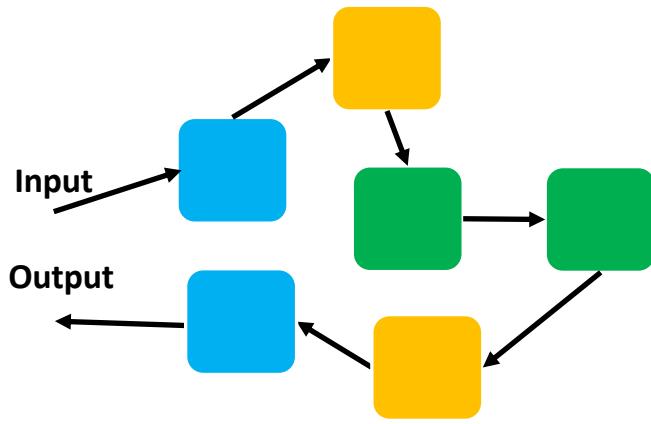
- Agents don't reference each other directly

Common ways of organizing agents

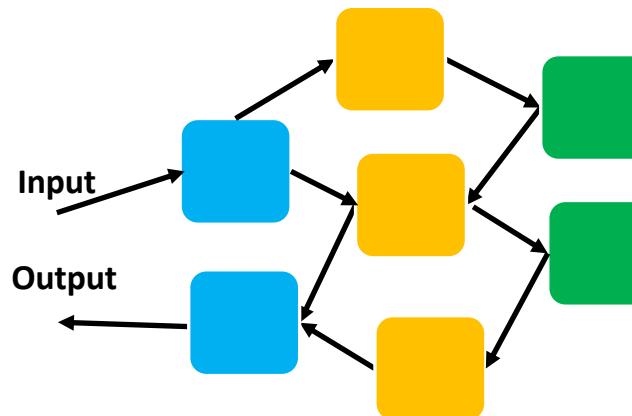
- **Worker agent** – Single agent does work in background
- **Layered network** – Agent uses agents from lower level
- **Pipeline processing** – Step-by-step processing

Comparison between Sequential, Task-based and Message passing programming

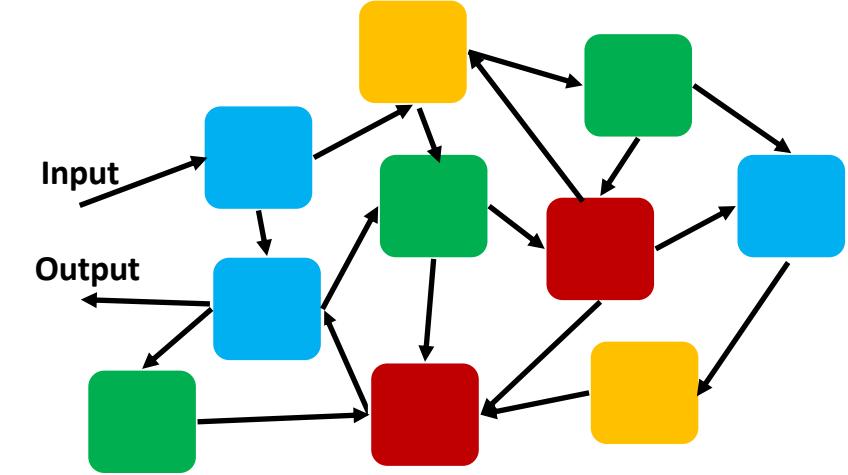
Sequential Programming



Task-Based Programming



Message-Passing Programming



Connecting agents

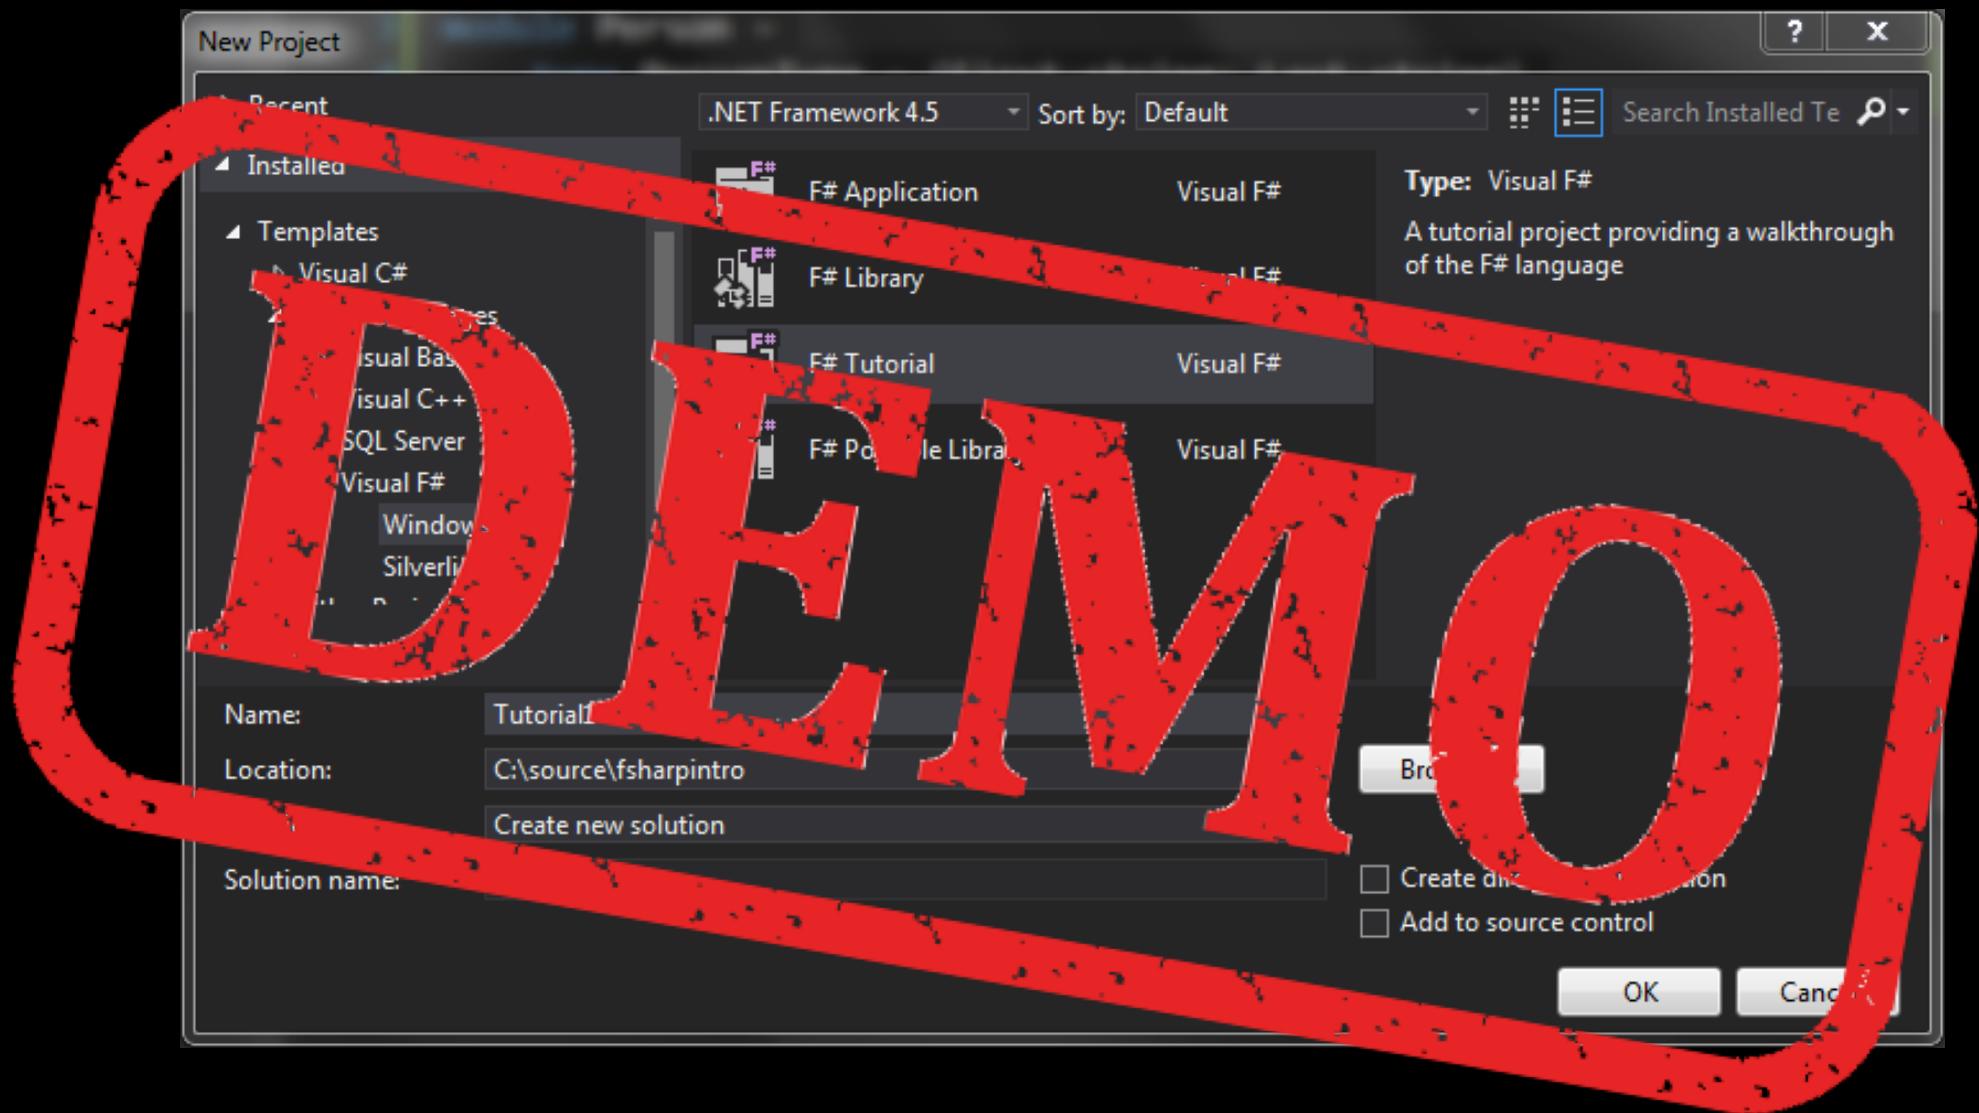
Applications consists of numerous agents

Agents communicate by sending messages

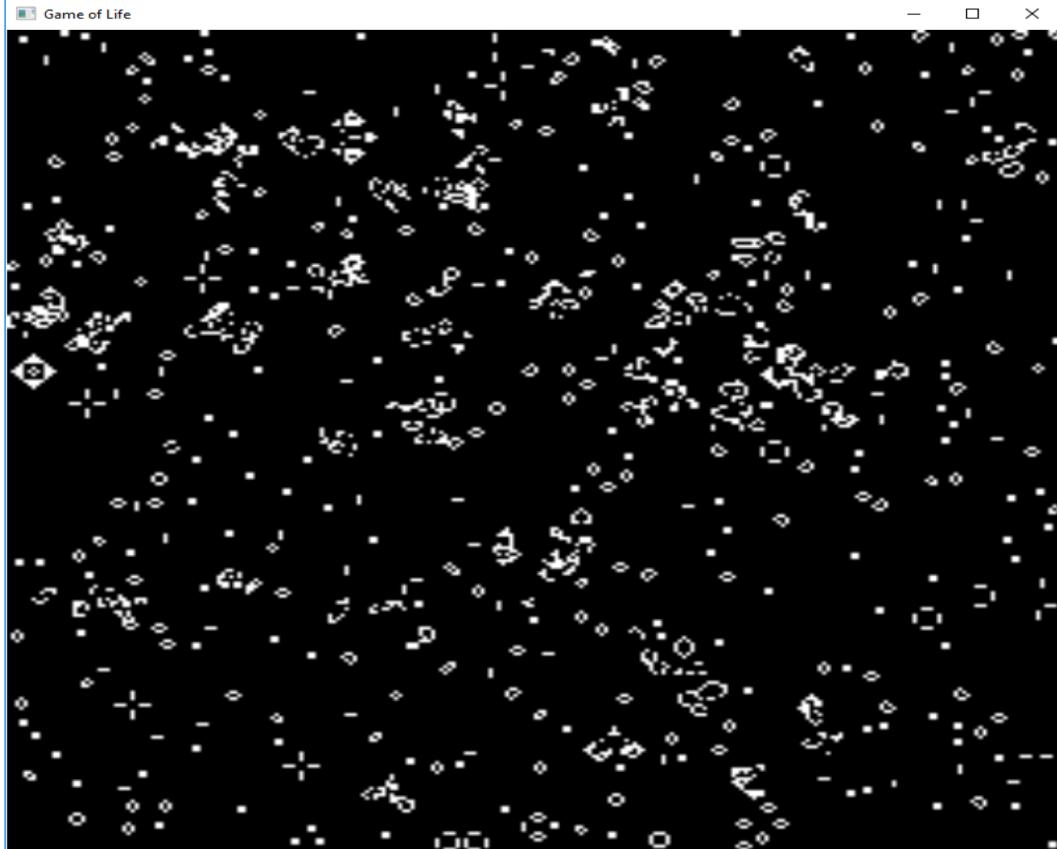
- Direct references are wrong!
- Difficult to follow, reuse and reconfigure

Decoupling of agents

- Expose events instead of sending messages
- Connect agents when creating them
- Use agent repository for dynamic configuration



Game of Life



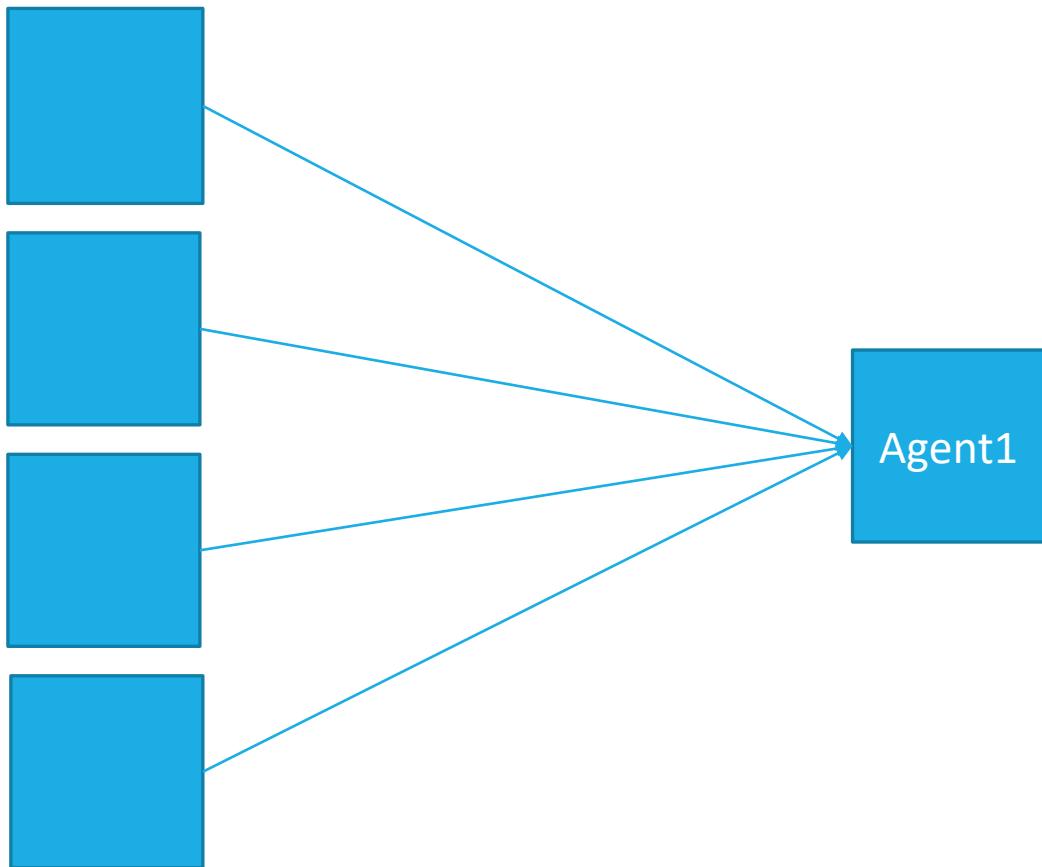
The Game of Life rules:

- Each cell with one or no neighbors dies, as if by solitude.
- Each cell with four or more neighbors dies, as if by overpopulation.
- Each cell with two or three neighbors survives.
- Each cell with three neighbors becomes populated (maximum).

Agent Patterns

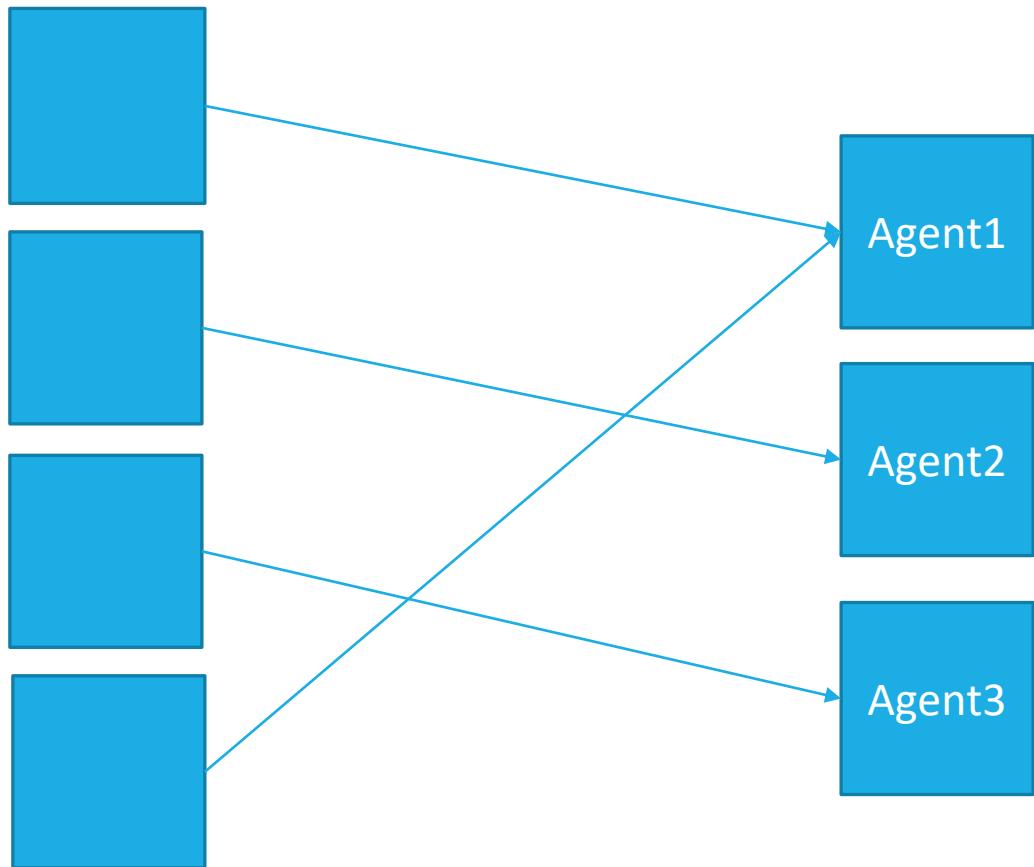
Agent patterns

Singleton



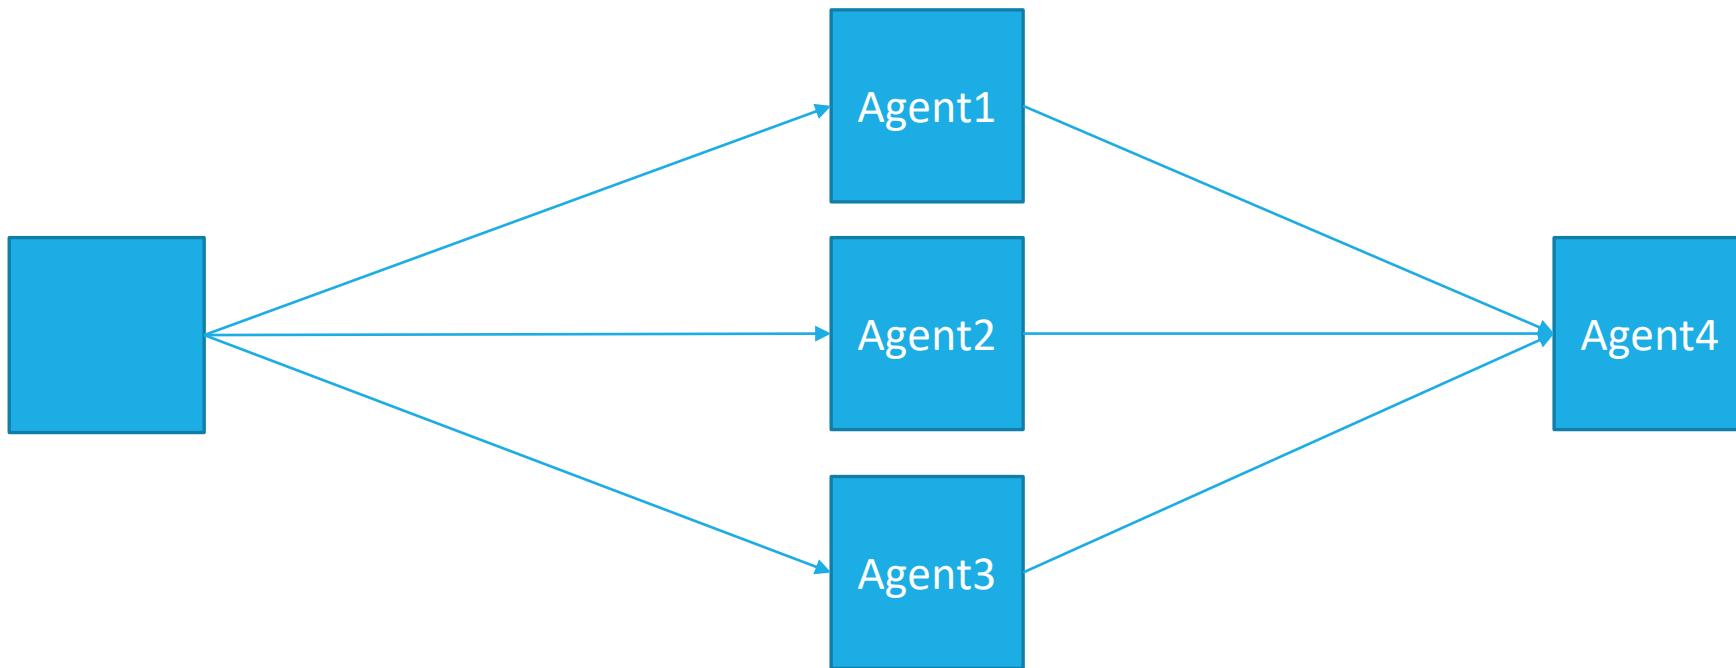
Agent patterns

Pool



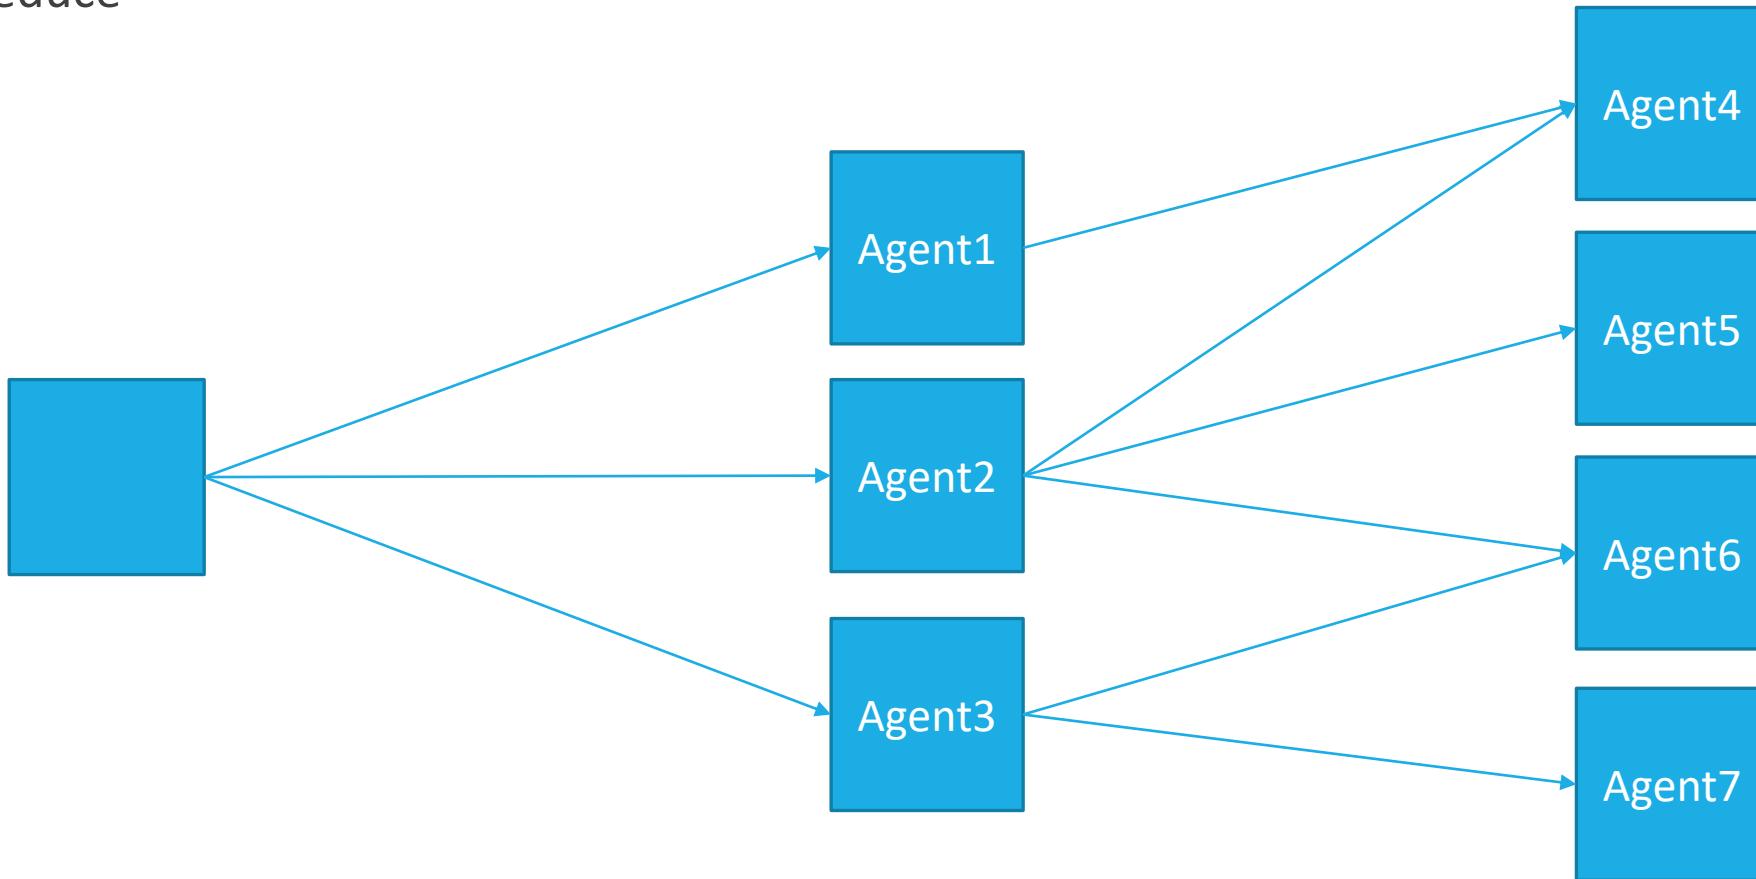
Agent patterns

Fork Join

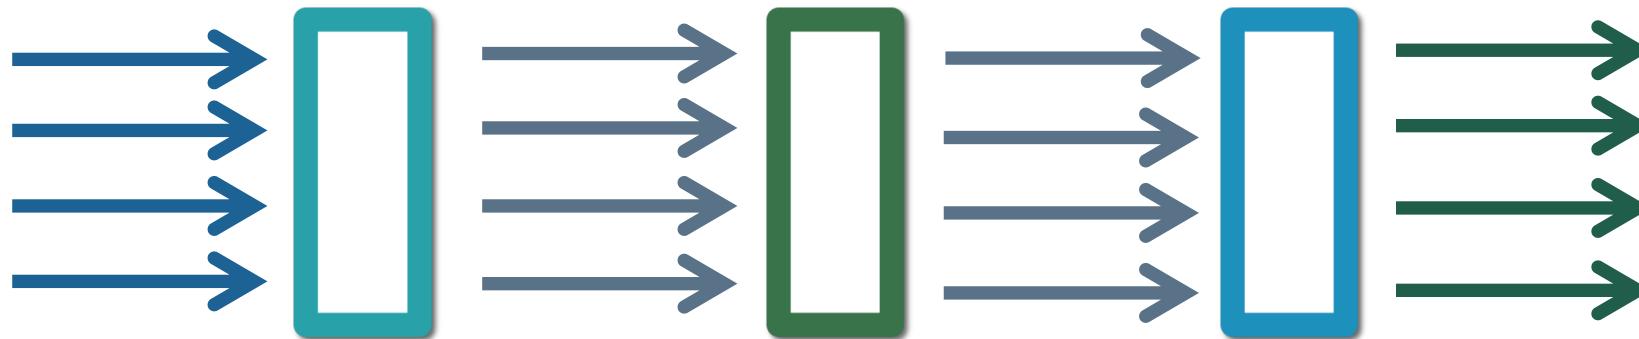


Agent patterns

Map Reduce



Pipeline Processing



- A pipeline is a set of data processing elements connected in series, so that the output of one element is the input of the next one. The elements of a pipeline are often executed in parallel or in time-sliced fashion; in that case, some amount of buffer storage is often inserted between elements

Agent Error Handling & Disposable

```
let errorAgent =
    Agent<int * System.Exception>.Start(fun inbox ->
        async { while true do
            let! (agentId, err) = inbox.Receive()
            printfn "an error '%s' occurred in agent %d" err.Message agentId })
```

```
let agent cancellationToken =
    new Agent<string>((fun inbox ->
        async { while true do
            let! msg = inbox.Receive()
            failwith "fail!" }), cancellationToken.Token)
    agent.Error.Add(fun error -> errorAgent.Post(error))
    agent.Start()
    agent
```

```
// (agent :> IDisposable).Dispose()
```

Agent Replying to the sender

Message carries input and a callback

```
type Message = string * AsyncReplyChannel<string>
```

Reply using the callback object

```
let echo = Agent<Message>.Start(fun agent ->
    async { while true do
        let! name, rchan = agent.Receive()
        rchan.Reply("Hello " + name) })
```

Asynchronous communication

```
let! s = echo.PostAndAsyncReply(fun ch -> "F#", ch)
```

Agent State Machine

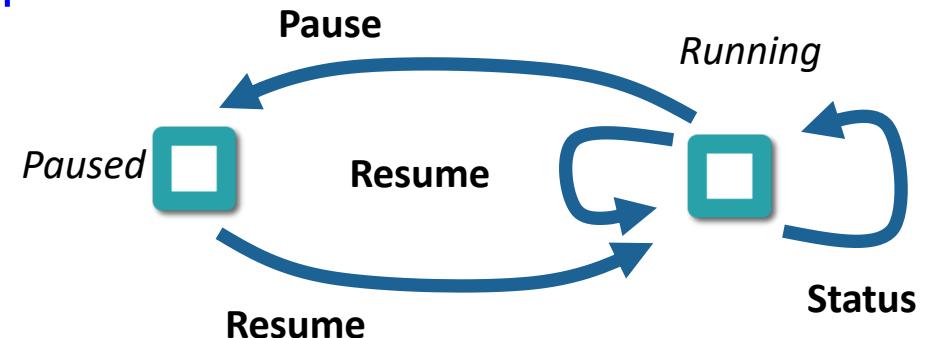
Multi-state agents use state machines

- Easy to implement as recursive functions
- Some states may leave messages in the queue

State transitions

- Accepting all messages, Asynchronously Receive and use pattern matching
- Waiting for a specific message, other messages stay in the queue

```
Agent.Start(fun agent ->
    let rec paused = agent.Scan (function
        | Resume -> Some (async {
            printfn "Resumed!"
            return! running })
        | _ -> None)
    and running = (* ... *) )
```



Agent Supervisors

```
let errorAgent =
    Agent<int * System.Exception>.Start(fun inbox ->
        async { while true do
            let! (agentId, err) = inbox.Receive()
            printfn "an error '%s' occurred in agent %d" err.Message agentId })
```

```
let agents10000 =
    [ for agentId in 0 .. 10000 ->
        let agent =
            new Agent<string>(fun inbox ->
                async { while true do
                    let! msg = inbox.Receive()
                    if msg.Contains("agent 99") then
                        failwith "fail!" })
        agent.Error.Add(fun error -> errorAgent.Post (agentId,error))
        agent.Start()
        (agentId, agent) ]
```

Scaling agents on demand

```
let urlList = [ ("Microsoft.com", "http://www.microsoft.com/");
               ("MSDN", "http://msdn.microsoft.com/");
               ("Google", "http://www.google.com") ]

let processingAgent() = Agent<string * string>.Start(fun inbox ->
    async { while true do
            let! name,url = inbox.Receive()
            let uri = new System.Uri(url)
            let webClient = new WebClient()
            let! html = webClient.AsyncDownloadString(uri)
            printfn "Read %d characters for %s" html.Length name } )

let scalingAgent : Agent<(string * string) list> = Agent.Start(fun inbox ->
    async { while true do
            let! msg = inbox.Receive()
            msg
            |> List.iter (fun x ->
                let newAgent = processingAgent()
                newAgent.Post x ))
```

Encapsulating agents

Create type and expose messages as members

```
type internal OnePlaceMessage<'T> =
| Put of 'T
| Get of AsyncReplyChannel<'T>

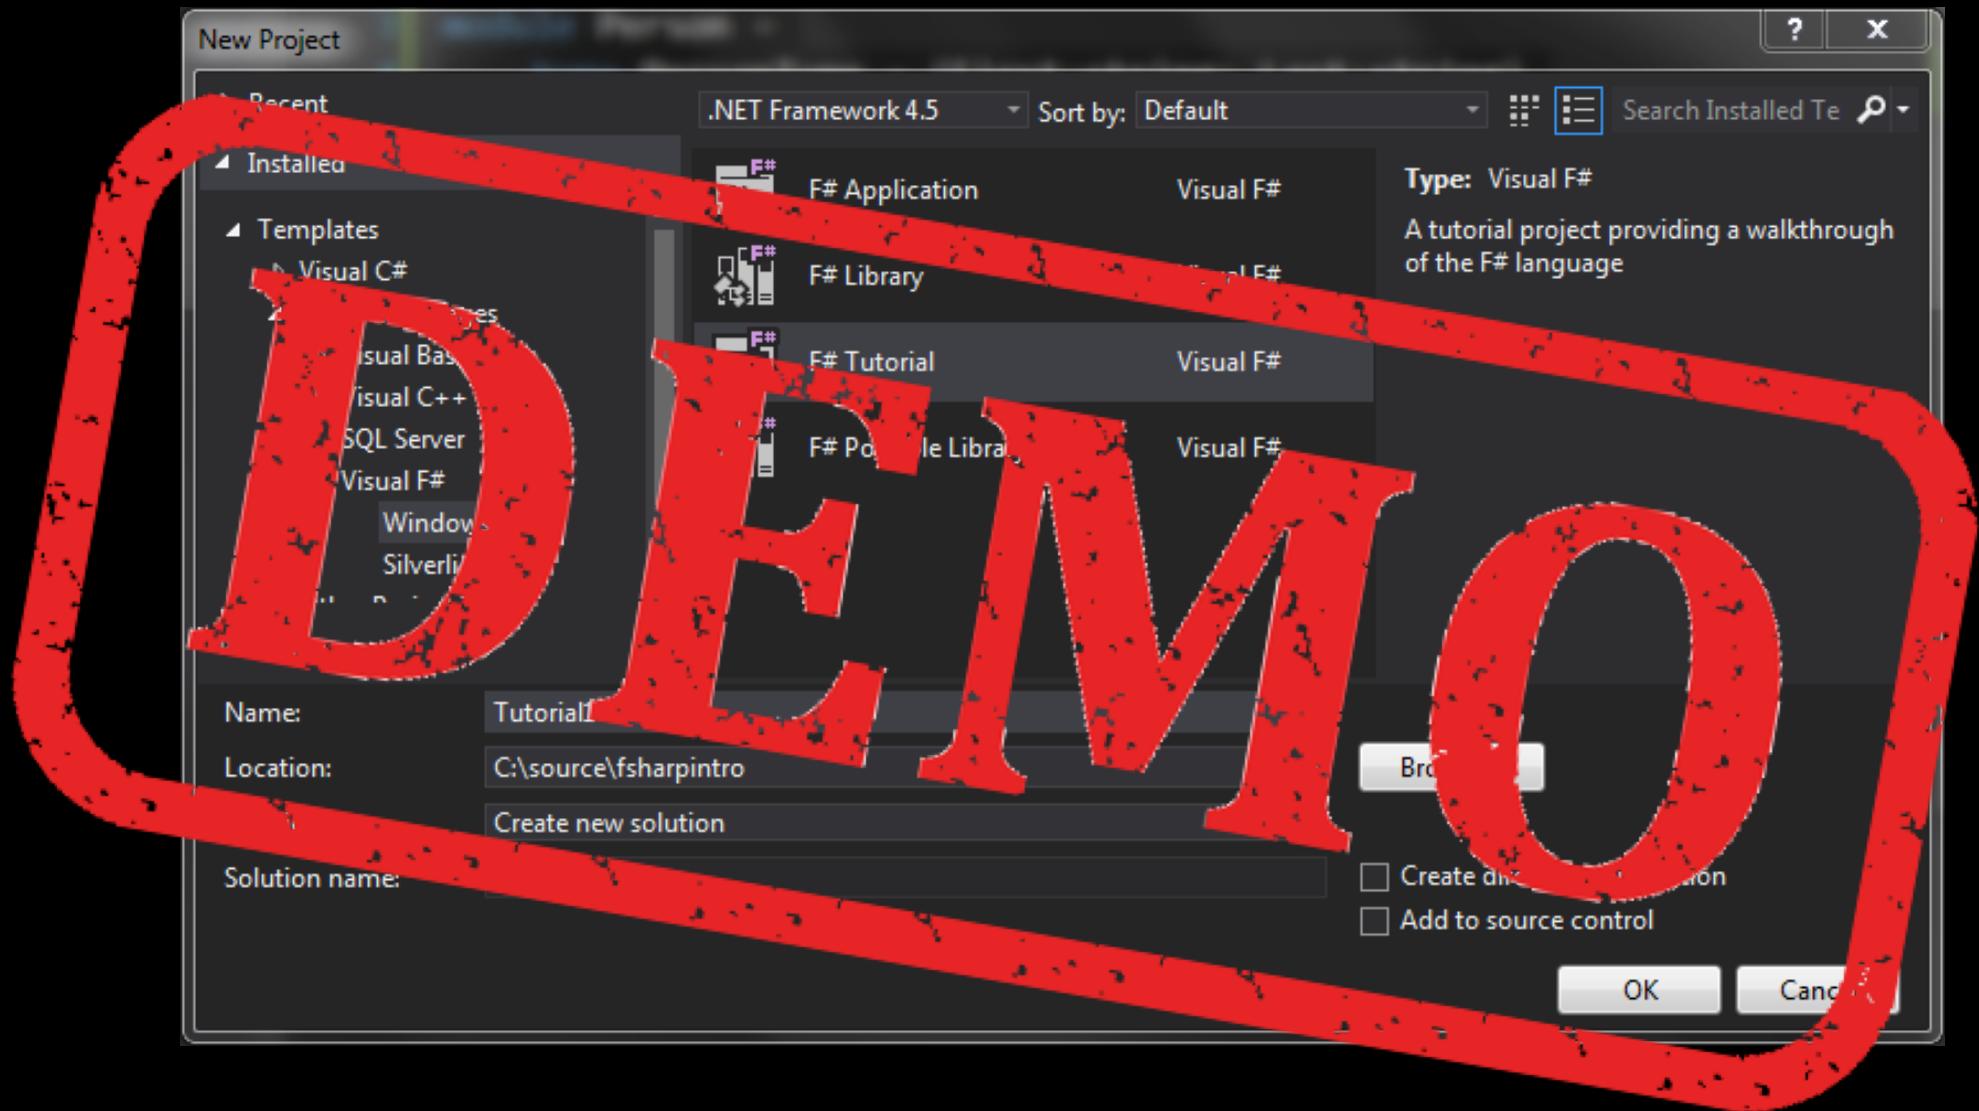
type OnePlaceAgent<'T>() =
    let agent = Agent.Start(fun agent -> (* ... *))
    member x.Put(value) = agent.Post(Put value)
    member x.AsyncGet() = agent.PostAndAsyncReply(Get)
```

- Ordinary methods for encapsulating Post
- Asynchronous methods when waiting for a reply

MailboxProcessor interoperability with C#

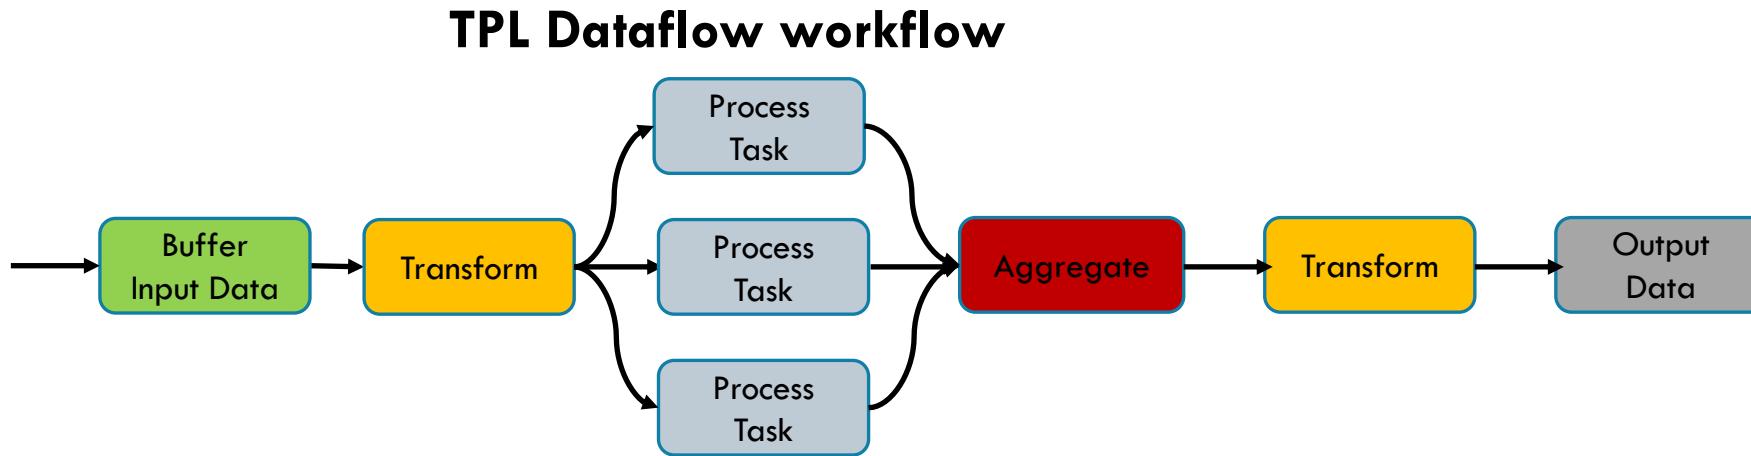
```
type internal OnePlaceMessage<'T> =
| Put of 'T
| Get of AsyncReplyChannel<'T>

type OnePlaceAgent<'T>() =
    let agent = Agent.Start(fun agent -> (* ... *))
    member x.Put(value) = agent.Post(Put value)
    member x.AsyncGet() = agent.PostAndAsyncReply(Get)
                           |> Async.StartAsTask
```

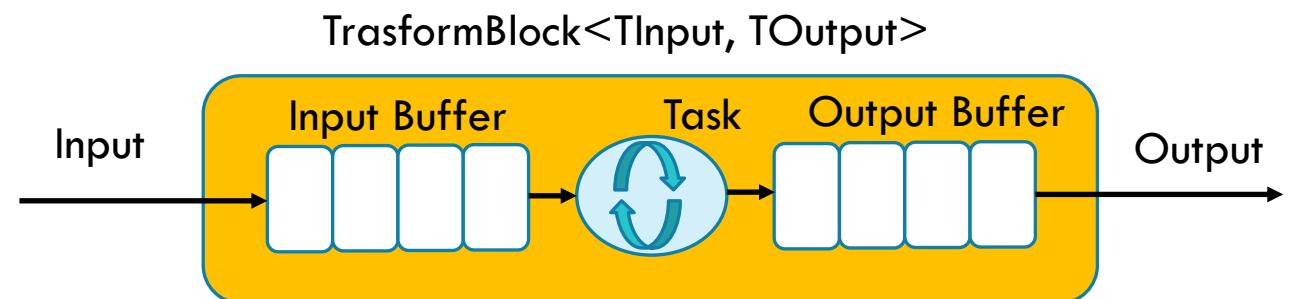
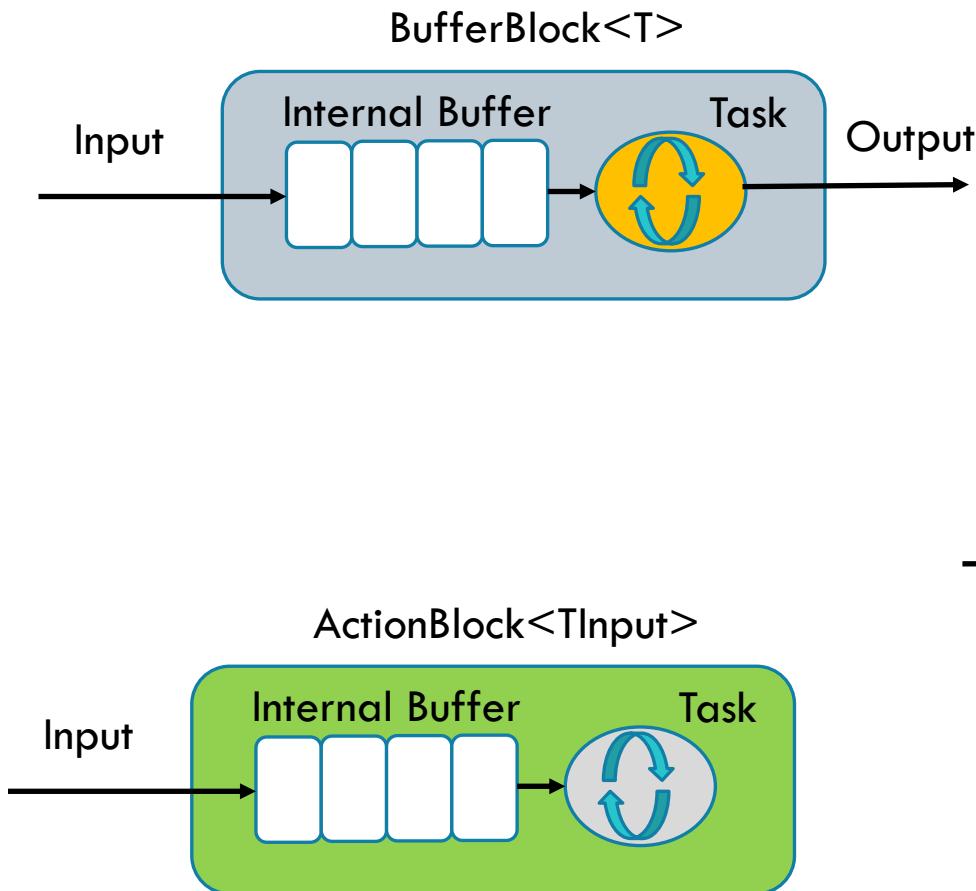


Agent in .NET and C#

TPL DataFlow blocks –design to compose



Some Dataflow blocks



Simple producer-consumer

```
BufferBlock<int> buffer = new BufferBlock<int>();  
  
async Task Producer(IEnumerable<int> values) {  
    foreach (var value in values)  
        buffer.Post(value);  
    buffer.Complete();  
}  
async Task Consumer(Action<int> process) {  
    while (await buffer.OutputAvailableAsync())  
        process(await buffer.ReceiveAsync());  
}  
async Task Run() {  
    IEnumerable<int> range = Enumerable.Range(0, 100);  
    await Task.WhenAll(Producer(range), Consumer(n =>  
        Console.WriteLine($"value {n}")));  
}
```

TPL DataFlow and Rx

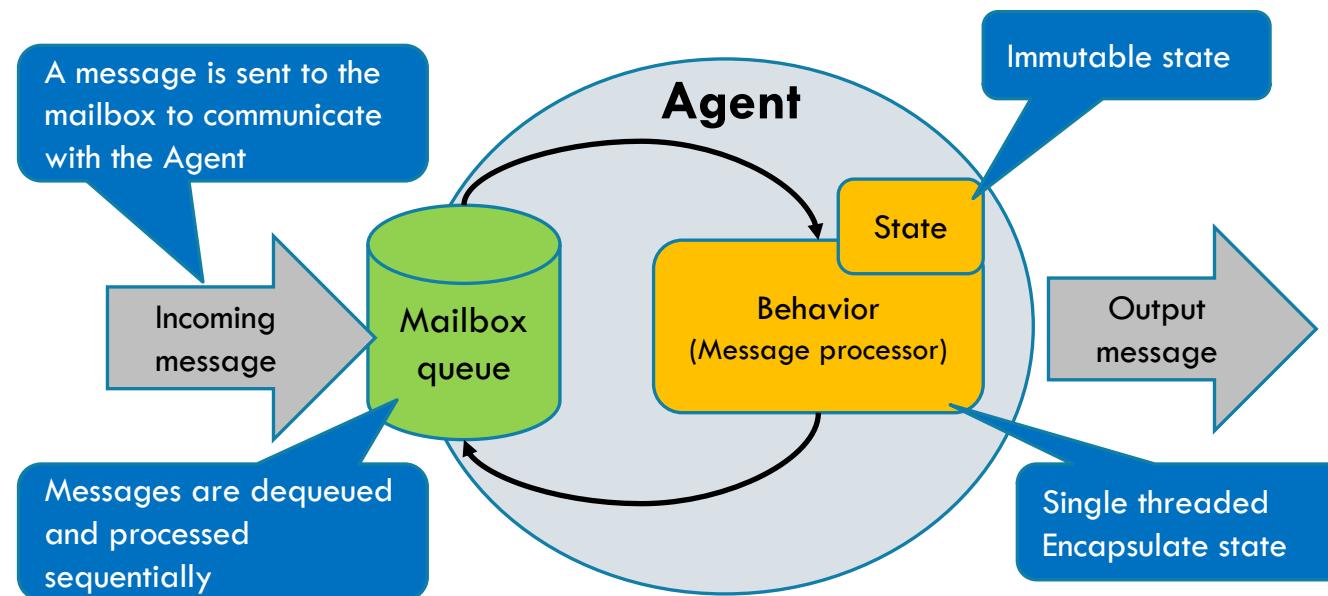
```
inputBuffer.LinkTo(compressor, linkOptions);
compressor.LinkTo(encryptor, linkOptions);

encryptor.AsObservable()

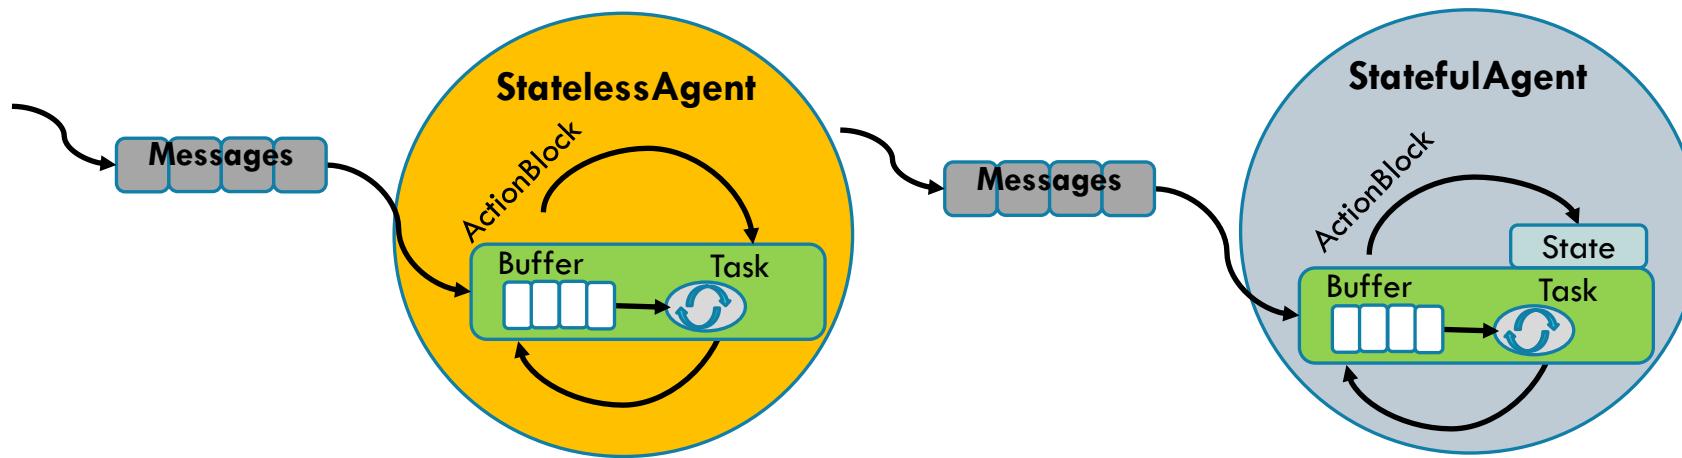
    .Scan((new Dictionary<int, EncryptDetails>(), 0),
(state, msg) => Observable.FromAsync(async() => {
    Dictionary<int,EncryptDetails> details, int lastIndexProc) = state;
    details.Add(msg.Sequence, msg);

    return (details, lastIndexProc);
}) .SingleAsync())
.SubscribeOn(TaskPoolScheduler.Default).Subscribe();
```

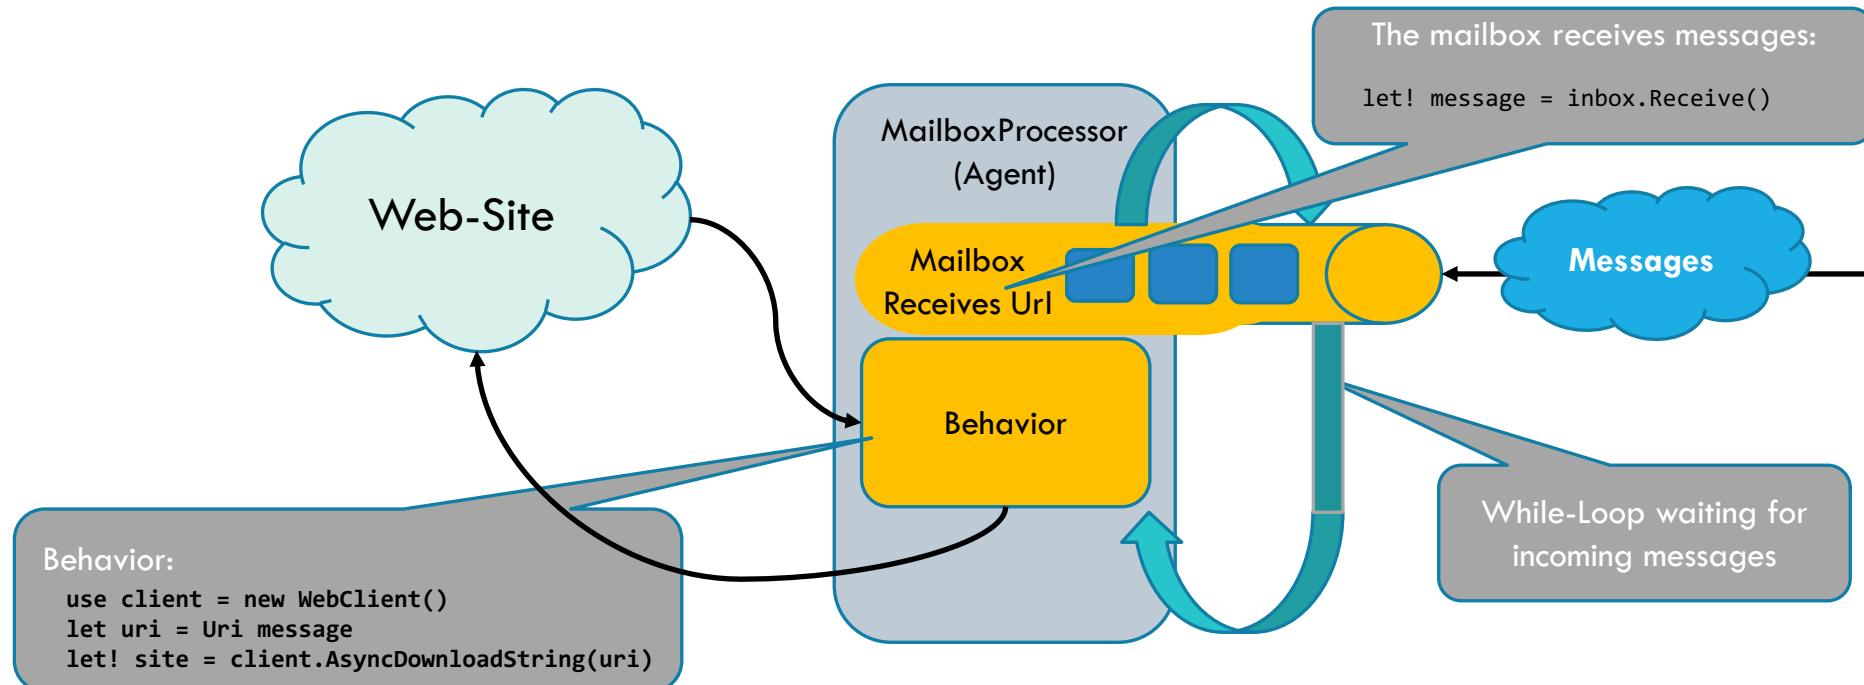
Agent anatomy



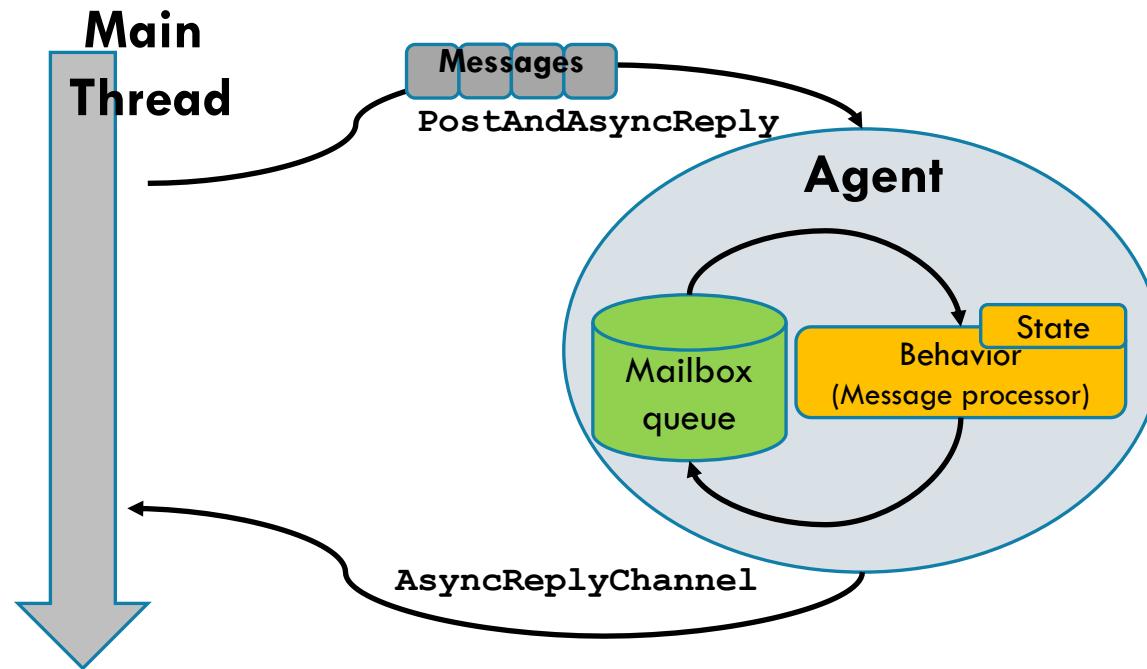
TPL DataFlow as Agent

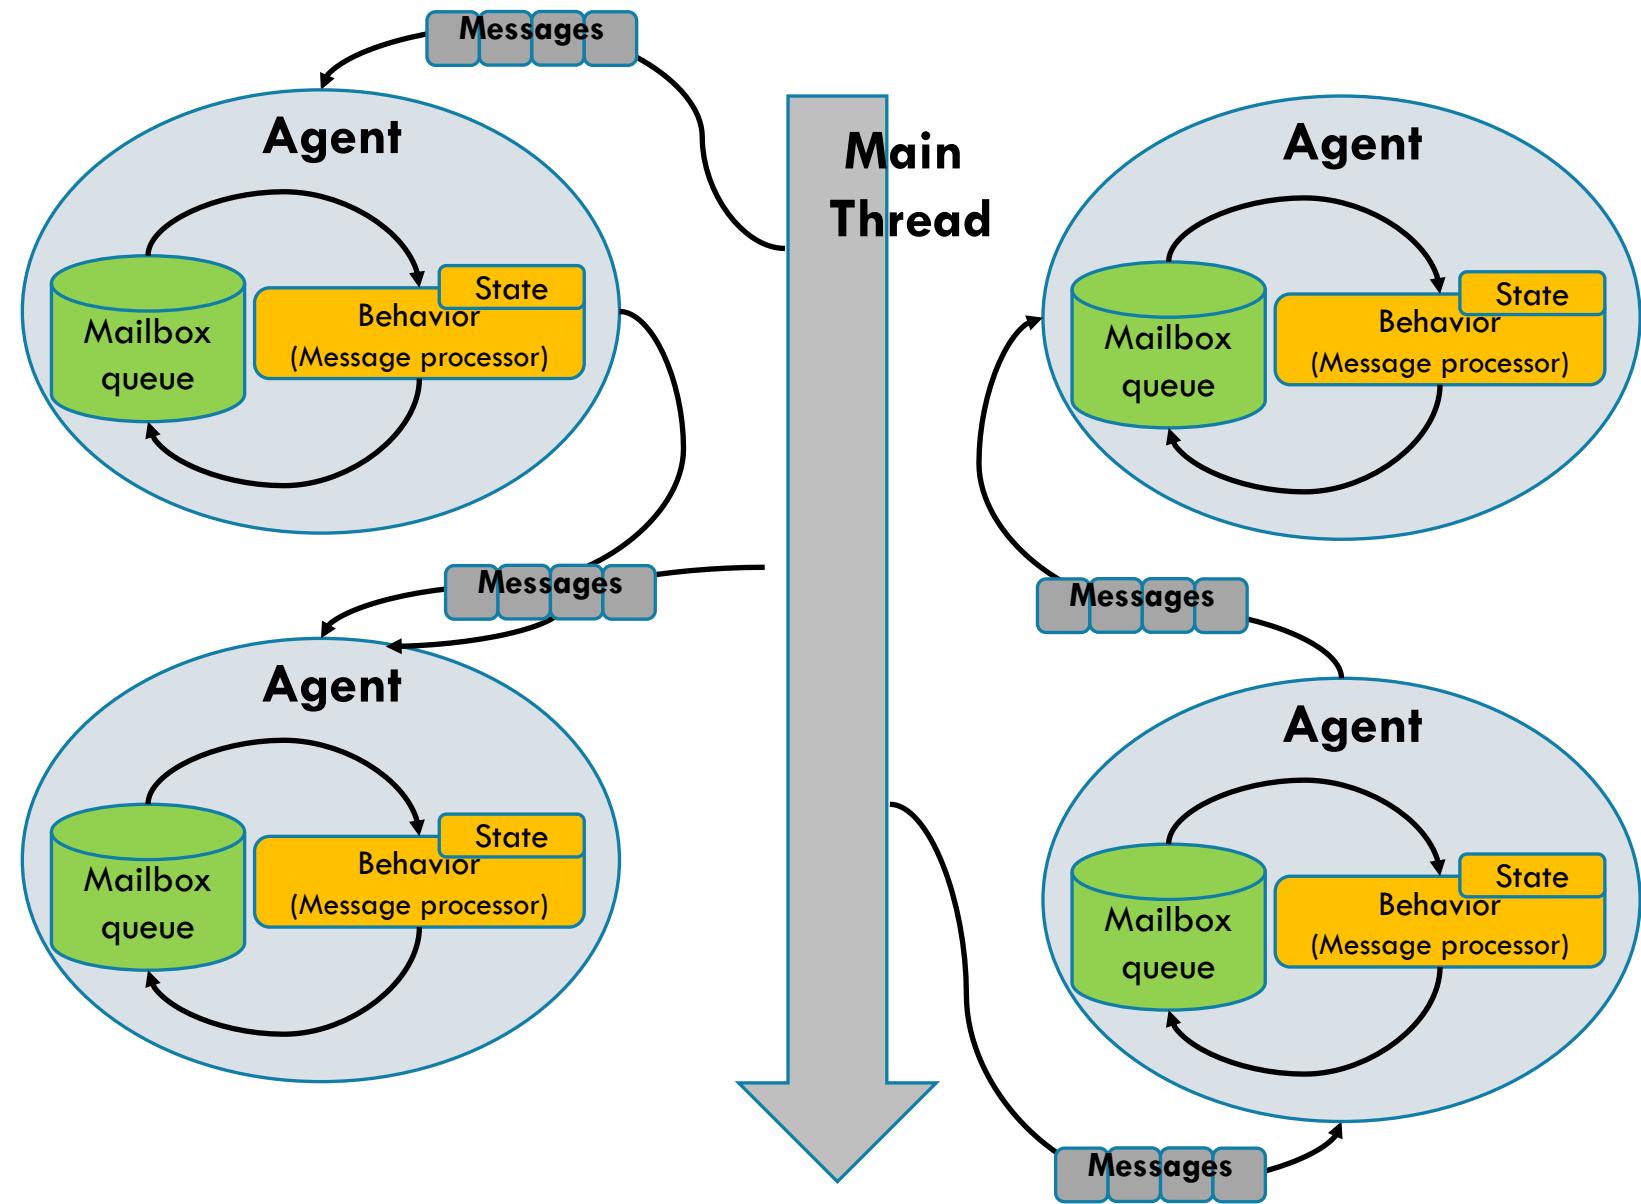


Agent anatomy



Agent two-way communication





Lab :

Implement a state full agent using
TPL DataFlow

Agent in C#

```
class StatefulDataflowAgent<TState, TMessage> : IAgent<TMessage>
{
    private TState state;
    private readonly ActionBlock<TMessage> actionBlock;

    public StatefulDataflowAgent(
        TState initialState,
        Func<TState, TMessage, Task<TState>> action,
        CancellationTokenSource cts = null)
    {
        state = initialState;
        var options = new ExecutionDataflowBlockOptions
        {
            CancellationToken = cts != null ?
                cts.Token : CancellationToken.None
        };
        actionBlock = new ActionBlock<TMessage>(
            async msg => state = await action(state, msg), options);
    }

    public Task Send(TMessage message) => actionBlock.SendAsync(message);
    public void Post(TMessage message) => actionBlock.Post(message);
}
```

TPL DataFlow a statefull agent

```
class StatefulDataFlowAgent<TState, TMessage>
{
    private TState state;
    private readonly ActionBlock<TMessage> actionBlock;

    public StatefulDataFlowAgent(
        TState initialState,
        Func<TState, TMessage, Task<TState>> action,
        CancellationTokenSource cts = null)
    {
        state = initialState;
        var options = new ExecutionDataFlowBlockOptions {
            CancellationToken = cts != null ?
                cts.Token : CancellationToken.None };

        actionBlock = new ActionBlock<TMessage>(async msg =>
            state = await action(state, msg), options);
    }

    public Task Send(TMessage message) => actionBlock.SendAsync(message);
    public void Post(TMessage message) => actionBlock.Post(message);
}
```

Agent two-way communication in C#

```
public Task<TReply> Ask(TMessage message)
{
    var tcs = new TaskCompletionSource<TReply>();
    actionBlock.Post((message, Some(tcs)));
    return tcs.Task;
}

public StatefulReplyDataflowAgent(TState initialState,
Func<TState, TMessage, Task<TState>> projection,
Func<TState, TMessage, Task<(TState, TReply)>> ask,
CancellationTokenSource cts = null)
{
    state = initialState;
    actionBlock = new ActionBlock<(TMessage, Option<TaskCompletionSource<TReply>>)>(
        async message =>
    {
        (TMessage msg, Option<TaskCompletionSource<TReply>> replyOpt) = message;
        await replyOpt.Match(
            none: async () => state = await projection(state, msg),
            some: async reply => {
                (TState newState, TReply replyresult) = await ask(state, msg);
                state = newState;
                reply.SetResult(replyresult);
            });
    });
}
```

TPL DataFlow caching web-sites downloaded

```
List<string> urls = new List<string> {
    "http://www.google.com",
    "http://www.microsoft.com",
    "http://www.bing.com",
    "http://www.google.com"
};

var agentStateful = Agent.Start(ImmutableDictionary<string, string>.Empty,
    async (ImmutableDictionary<string, string> state, string url) => {
        if (!state.TryGetValue(url, out string content))
            using (var webClient = new WebClient())
                content = await webClient.DownloadStringTaskAsync(url);
        await File.WriteAllTextAsync(createFileNameFromUrl(url), content);
        return state.Add(url, content);
    }
    return state;
}) ;

urls.ForEach(url => agentStateful.Post(url));
```

Agent fold-over state and messages (Aggregate)

```
Agent<ImmutableDictionary<string, string>, Empty> = agent =>
    agent.Add(url, content) => {
        if (!state.TryGetValue(url, out string content))
            using (var webClient = new WebClient())
            {
                content = await webClient.DownloadStringTaskAsync(url);
                await File.WriteAllTextAsync(createFileNameFromUrl(url), content);
            }
        return state.Add(url, content);
    };
};
```

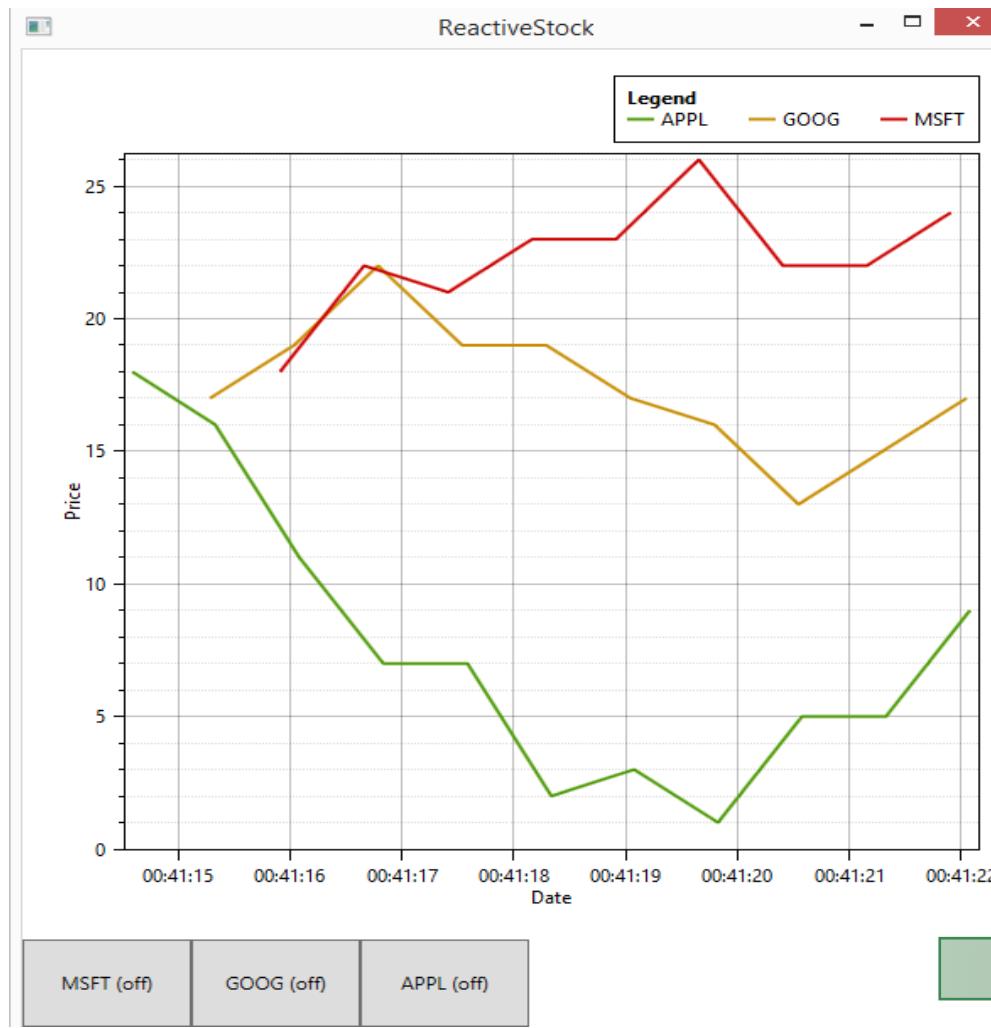
Agent fold-over state and messages (Aggregate)

```
urls.Aggregate(ImmutableDictionary<string, string>.Empty,  
    async (state, url) => {  
    if (!state.TryGetValue(url, out string content))  
        using (var webClient = new WebClient())  
    {  
        content = await webClient.DownloadStringTaskAsync(url);  
        await File.WriteAllTextAsync(createFileNameFromUrl(url), content);  
        return state.Add(url, content);  
    }  
    return state;  
});
```

Lab :

Agent Stock Ticker

Agent Stock Ticker



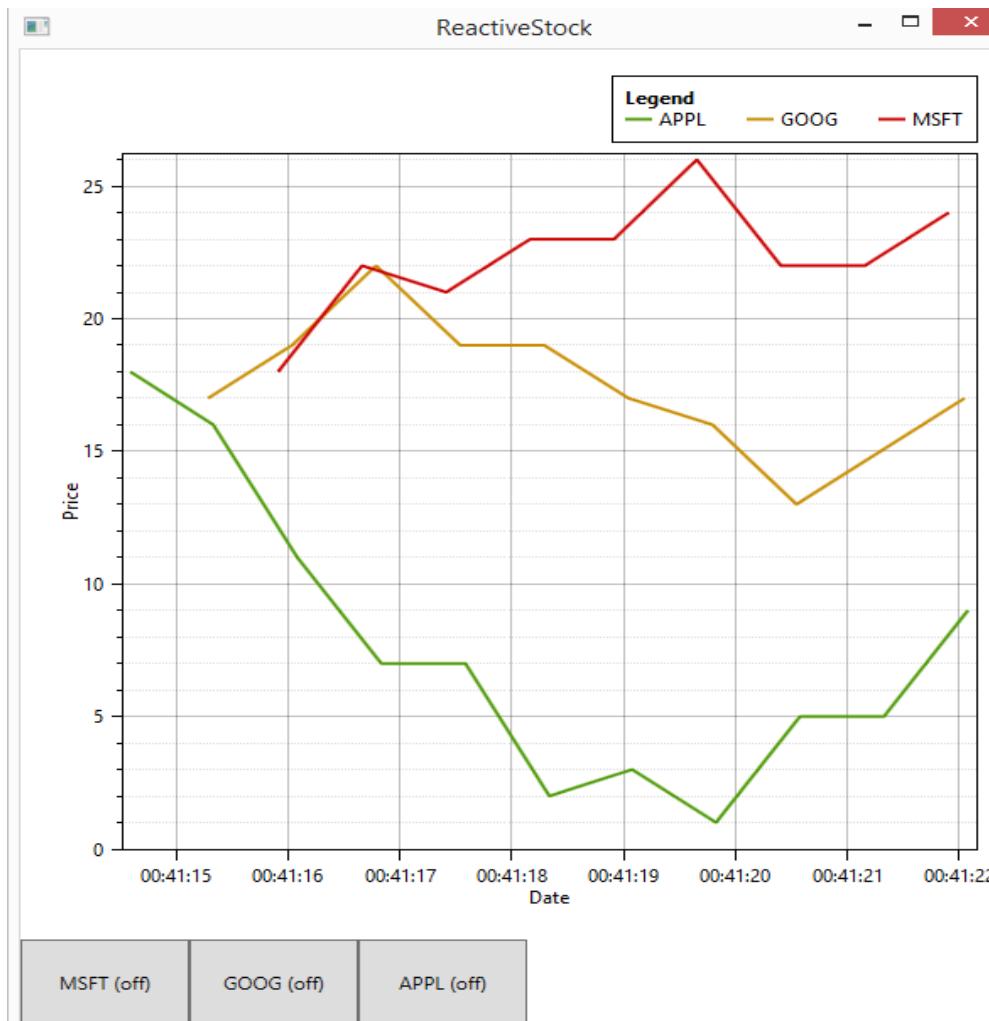
Agent
Charting

Agent
Coordinator

Agent Stock

Update
Stock Prices

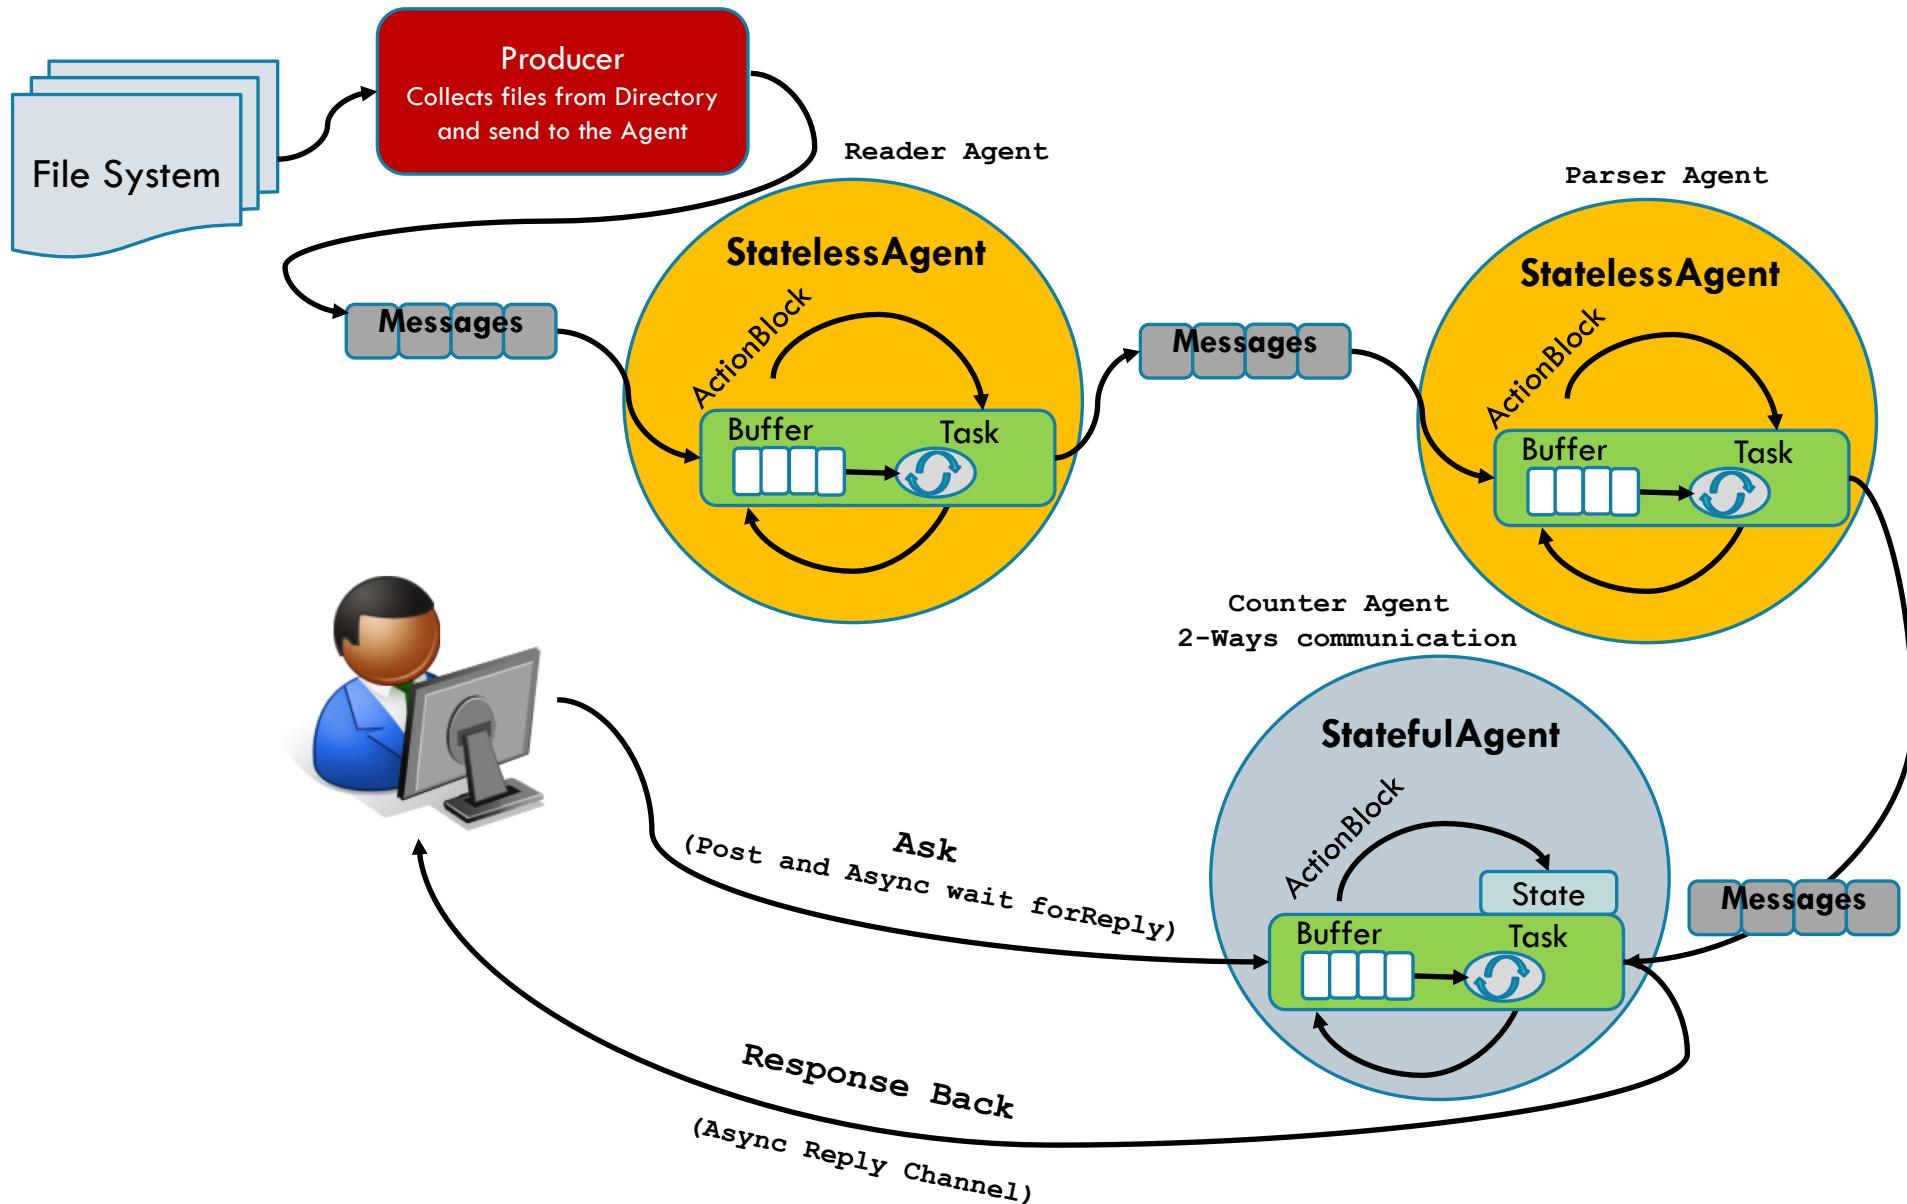
Agent Stock Ticker



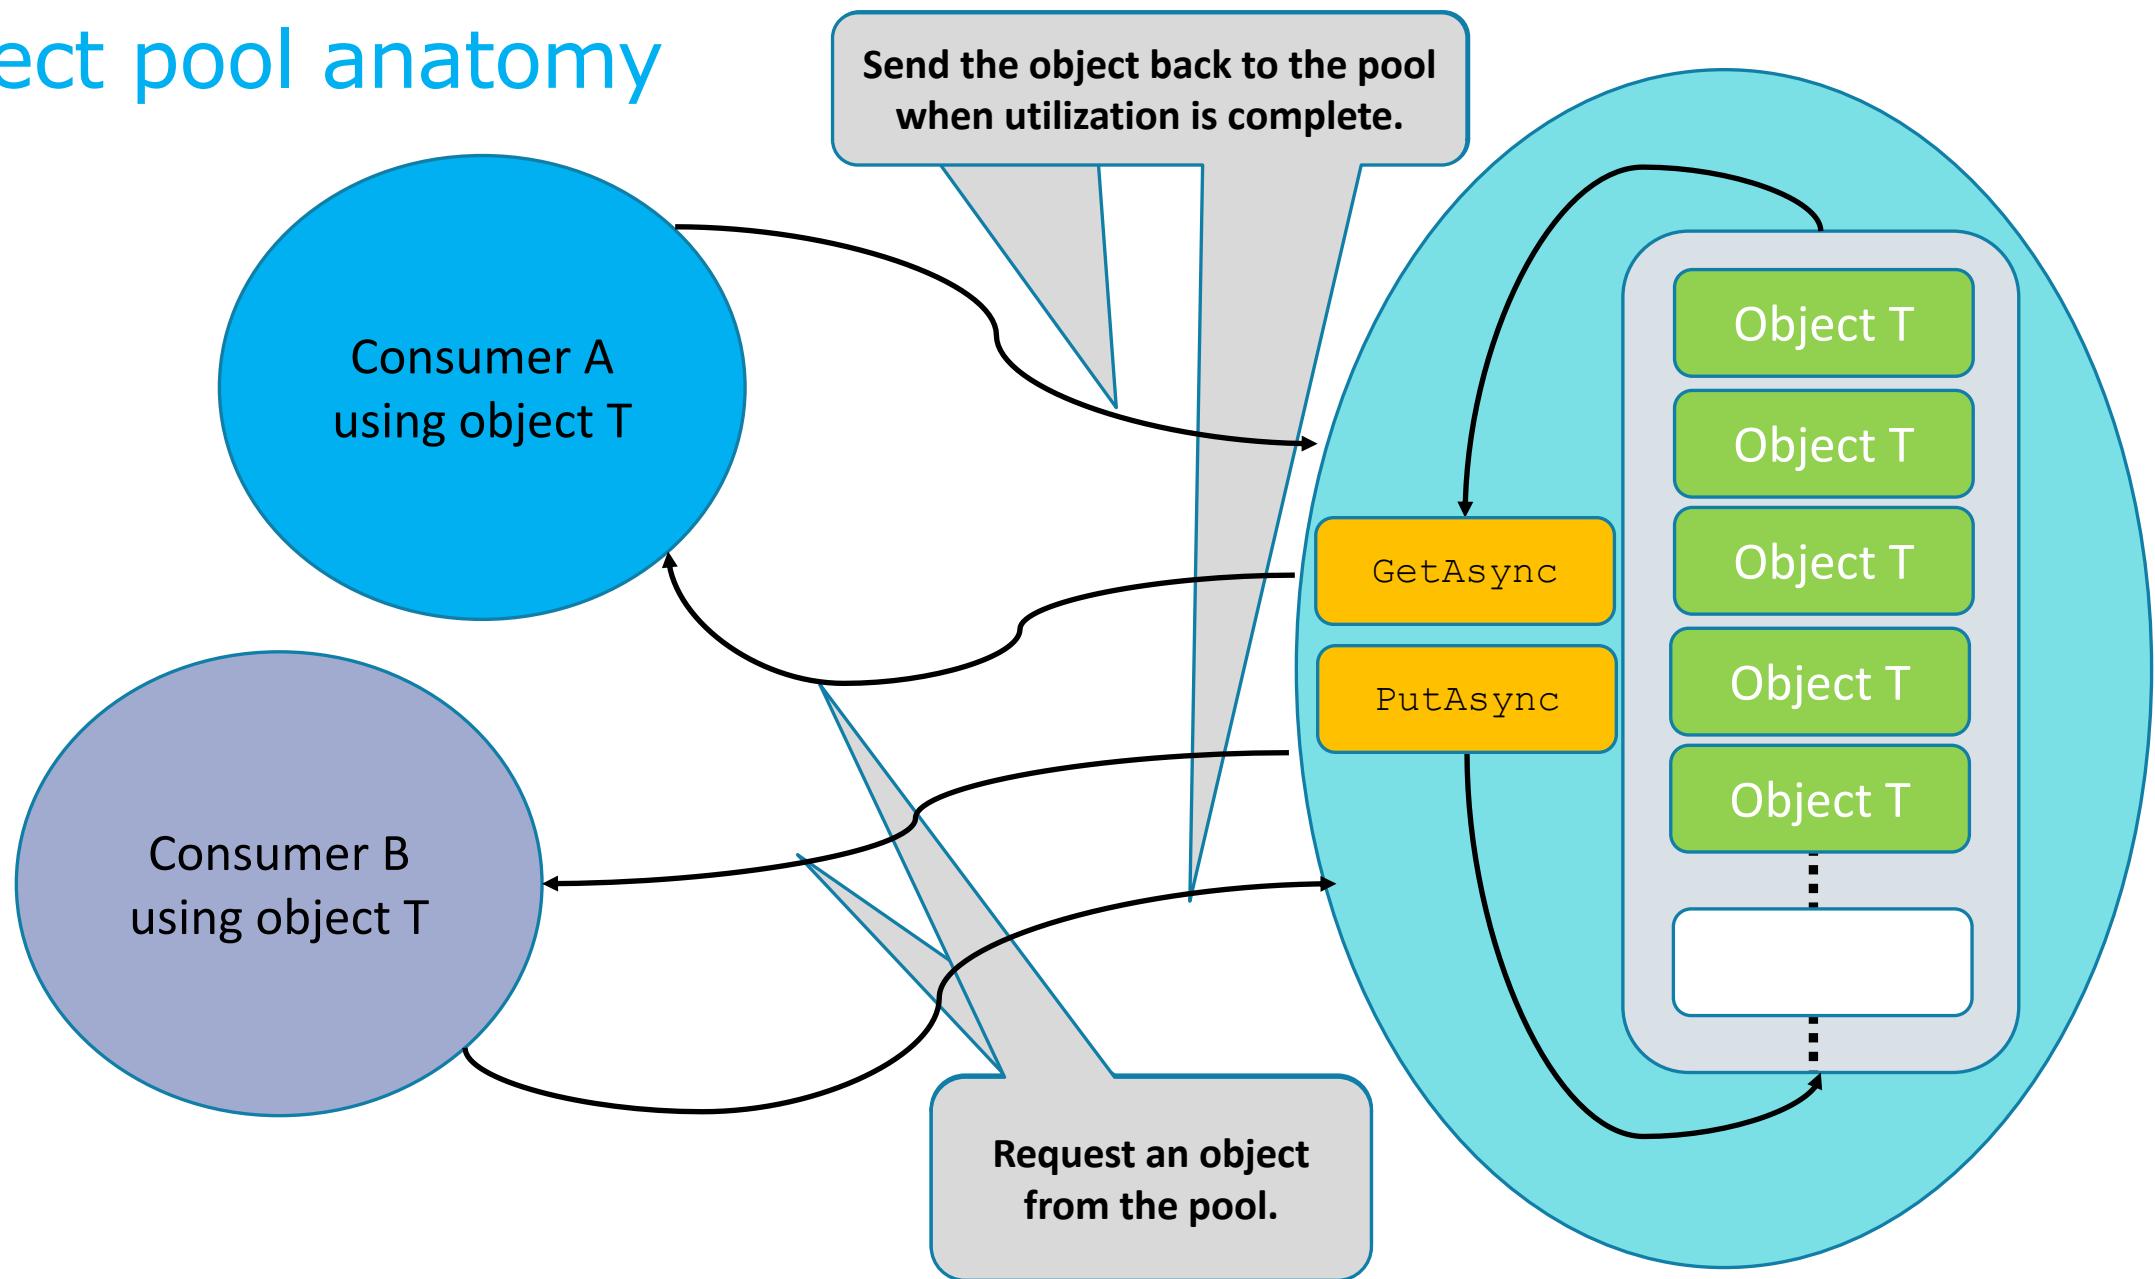
Tasks:

1. Create Agent hierarchy Children-Parent (sub-pub)
 - Create Agent Stock (one per stock symbol)
 - Create Agent Coordinator for subscribe/unsubscribe Agent Stocks
2. Connect Agents using messages

Agent for producer and consumer pattern



Object pool anatomy



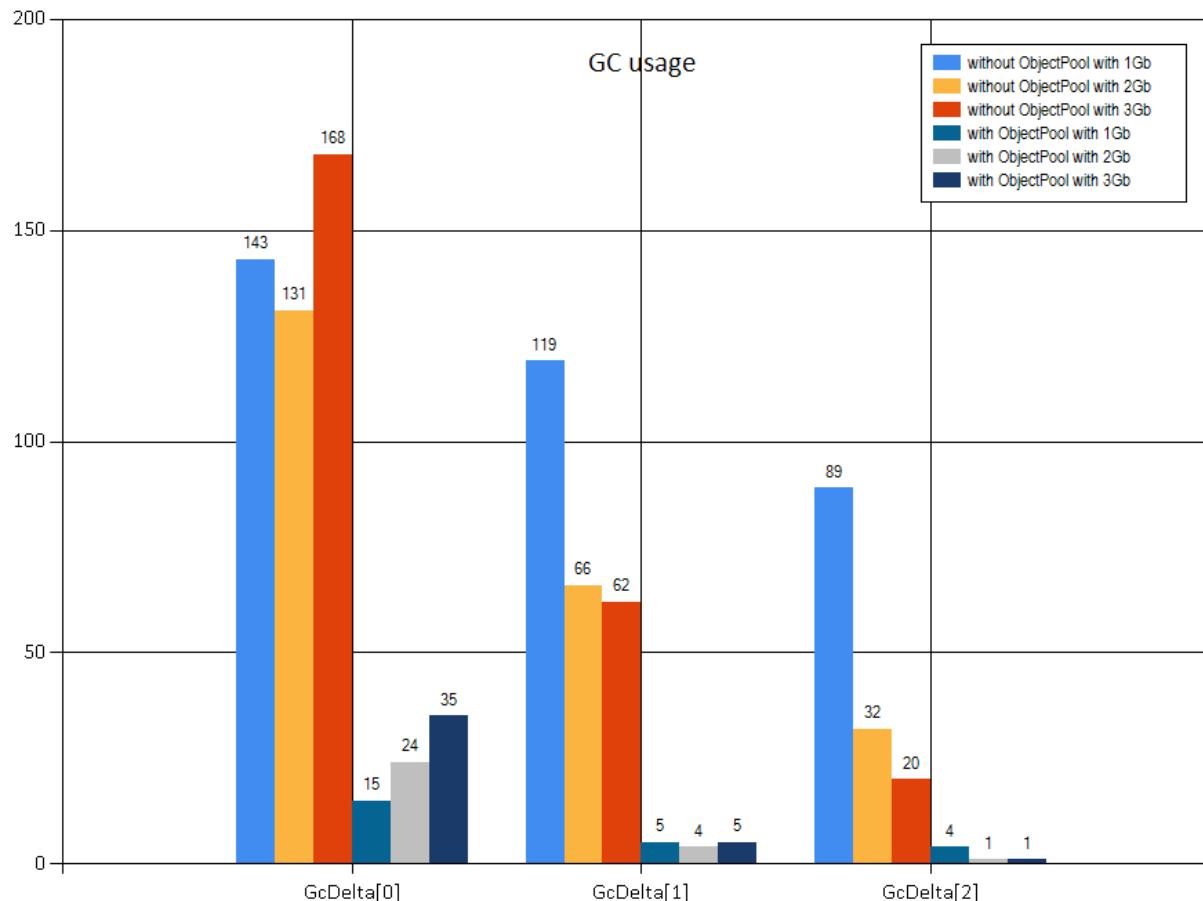
Async-PoolObject

```
public class ObjectPoolAsync<T>
{
    private readonly BufferBlock<T> buffer;
    private readonly Func<T> factory;
    private readonly int msecTimeout;
    private int currentSize;

    public ObjectPoolAsync(int initialCount, Func<T> factory, CancellationToken cts, int msecTimeout = 0)
    {
        this.msecTimeout = msecTimeout;
        buffer = new BufferBlock<T>(~#A
            new DataflowBlockOptions { CancellationToken = cts });
        this.factory = () => {
            Interlocked.Increment(ref currentSize);
            return factory();
        };
        for (int i = 0; i < initialCount; i++)
            buffer.Post(this.factory()); ~#B
    }
    public int Size => currentSize;

    public Task<bool> PushAsync(T item) => ...
    public Task<T> GetAsync(int timeout = 0) ...
```

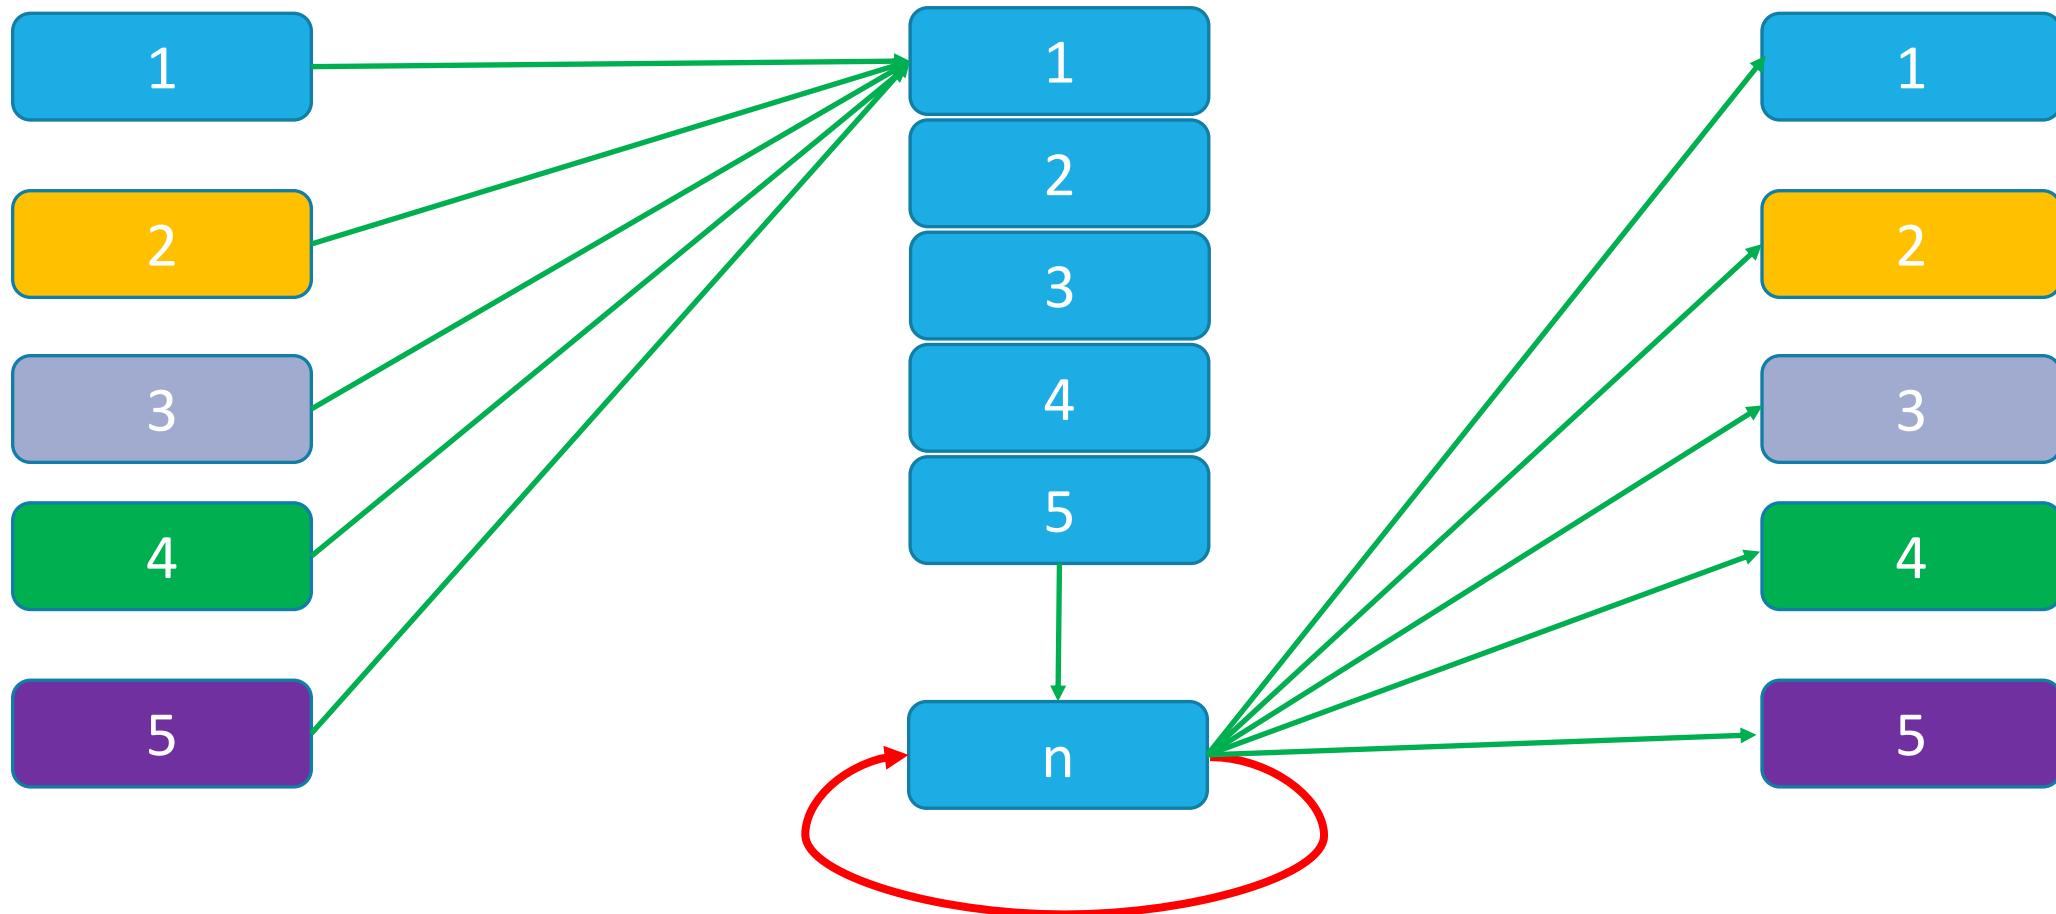
Async-PoolObject benefit



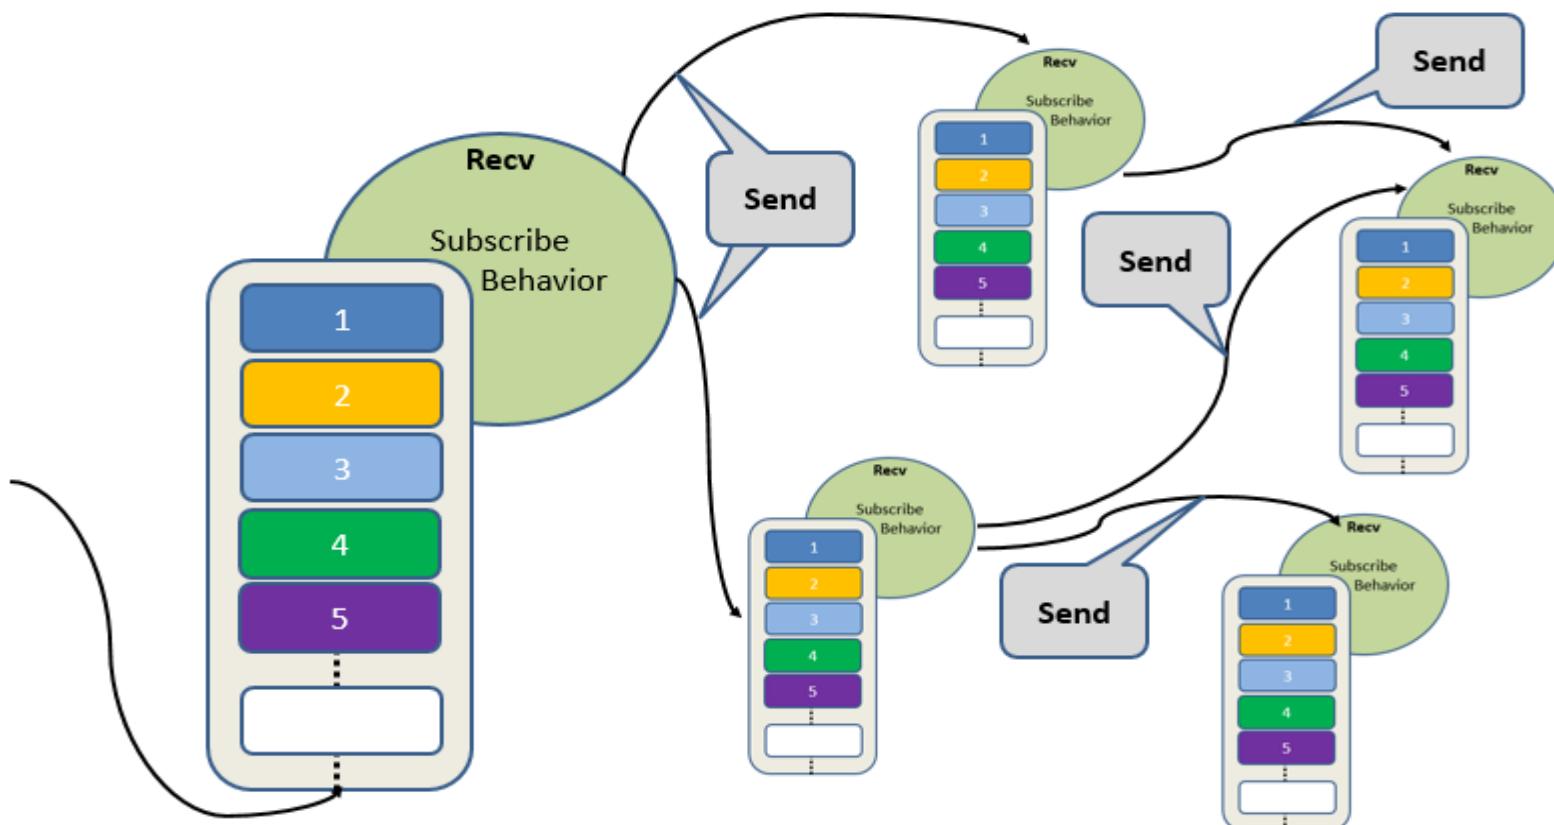
Channel



Channel



Channel



```
type [<Sealed>] Channel<'T>() =
    let channelAgent =
        MailboxProcessor<ChannelMsg<'T>>.Start(fun inbox ->
            let readers = Queue()
            let writers = Queue()
            let rec loop () = async {
                let! msg = inbox.Receive()

                match msg with
                | Read ok when writers.Count = 0 ->
                    readers.Enqueue ok
                    return! loop()
                | Read ok -> let value, cont = writers.Dequeue()
                    Pool.Spawn cont
                    ok value
                    return! loop()
                | Write (x,ok) when readers.Count = 0 ->
                    writers.Enqueue(x, ok)
                    return! loop()
                | Write (x,ok) -> let cont = readers.Dequeue()
                    Pool.Spawn ok
                    cont x
                    return! loop() }

            loop() )

    member this.Read(read) = channelAgent.Post (Read read)
    member this.Write(v, ok) = channelAgent.Post (Write (v,ok))

    member inline this.Read() = Async.FromContinuations(fun (ok, _, _) -> this.Read ok)
    member inline this.Write x = Async.FromContinuations(fun (ok, _, _) -> this.Write(x,ok))
```

Lab :

Channel in action (broadcasting)



The tools we use have a profound (and devious!) influence on our thinking habits, and, therefore, on our thinking abilities.

-- Edsger Dijkstra